

Abstract	2
Problem Statement	2
Goals	3
Non-Goals	3
Data Sources	4
Data Exploration	4
Exploratory Data Analysis (EDA)	5
Tools Used to Ingest Data	7
Security Checklist, Privacy, and Other Risks	8
Data Preparation	9
Feature Selection, Creation, and Transformation	10
Class Imbalance	11
Train, Test, Validation Splits	11
Model Training	11
Built-in Algorithm Approach	12
Data Training	12
Instance Size and Count	12
Model Evaluation	12
Future Enhancements	13
References	14
Appendix	15

Authors: Graham Ward and Tanya Ortega

Company Name: Pacific Freight Logistics

Company Industry: Aviation Shipping and Transportation Analytics

Company Size: 550 employees

Github Repository: <https://github.com/gw-00/ADS-508-Project>

Abstract

Shipping and logistics companies depend on efficient air cargo transportation to meet delivery deadlines and maintain supply chain stability. However, shipment delays and cancellations are often caused by unpredictable factors such as weather conditions, customs processing delays, and airport congestion. These disruptions result in financial losses, missed contracts, and challenges in logistics planning.

Problem Statement

The logistics industry suffers challenges due to air cargo delays and cancellations, leading to cargo not arriving on time, financial losses, and the implementation of last-minute contingency plans. These unforeseen challenges also impact downstream businesses that rely on timely deliveries for proper inventory management and customer satisfaction. As a shipping aviation and transportation company working with San Diego International Airport, Pacific Freight Logistics will provide predictive insights to help logistics managers anticipate delays and cancellations, allowing for better scheduling and a more reliable supply chain.

Goals

1. Improve flight delay predictions by developing a machine learning model that integrates weather data and historical flight performance, achieving at least **75%** accuracy in classifying flights as delayed/canceled or on time.
2. Reduce average cargo flight delay times by **2%** compared to historical benchmarks by providing logistics teams with delay probability scores that inform proactive decision-making.
3. Establish a correlation between weather conditions, flight information, and delays. We will perform a regression analysis to quantify the impact of weather conditions on delay, developing a coefficient to determine the impact of each weather variable on the probability of a delay. The goal would be to identify the greatest weather contributor that impacts flight delays and cancellations at San Diego International Airport.

Non-Goals

1. This project will not predict the length of time of a delay and will instead focus on just the classification of either delay/cancel and not.
2. The system will not make real-time operational decisions—it will serve as a decision-support tool for shipping companies.
3. We will not integrate direct air traffic control data and radar but will use available historical flight records and weather METAR.

4. We are not replacing existing transportation management systems (TMS) but enhancing them with predictive analytics to inform decision-making related to optimizing logistics routes for shipping companies.

Data Sources

Processed data that is ready for modeling and to be queried will be uploaded into Amazon S3 buckets and queried using AWS Athena.

1. NCEI ISD (National Center for Environmental Information Integrated Surface Database)

<https://www.ncei.noaa.gov/products/land-based-station/integrated-surface-database>

- a. Hourly and weather dependent observations taken for San Diego International Airport over the course of a two year period (2023 - 2025)

2. Bureau of Transportation Statistics *"On-Time: Reporting Carrier On-Time Performance"*

<https://www.transtats.bts.gov/Fields.asp?gnoyr VQ=FG>

- a. Data Collection of every single flight and its associated on time performance broken down annually for all airports in California. Later the data is filtered to only include flights departing San Diego International Airport


Data Exploration

This project began with two primary raw datasets as mentioned above, the historical flight performance and weather observations. Before uploading all the data into our raw data S3 bucket, the on time reports and weather reports were merged with their respective CSV files in chronological order and then uploaded individually as the raw versions to the raw S3 bucket. METAR weather reports had to be extracted and parsed into their separate categories using REGEX code to break down METAR code and ensure that each weather variable was extracted

into its own column for analysis.. Once the raw files were loaded into the S3 buckets they were pulled into a SageMaker notebook to clean and integrate with one another with the goal of each flight being associated with the legally defined weather for its take off time. Date and time formats were standardized in order to accurately match flights with their respective weather for the time block. Pandas was used to merge both the datasets into one for performing Exploratory Data Analysis (EDA).


Exploratory Data Analysis (EDA)

During the EDA phase of the project we ensured the flight and weather data were merged correctly and then explored questions regarding the distribution of on-time flights, cancelled flights and delayed flights. We also addressed missing values in the '*wind_dir_degrees*' column by setting all them to 360, which was the predominant wind direction. When there is a missing wind direction it typically means that winds are calm and variable with no predominant direction. With this in mind we still needed a way to retain as much flight information as possible and setting values that were missing to the predominant wind direction would be acceptable for our use cases. Figure 1, is a windrose, a common depiction of where the winds for a given location come from and how strong they are. The darker color segments point towards lower speeds but each concentric ring indicates frequency of the wind direction. What we see from this windrose is that predominantly there are light winds out of the North for San Diego International Airport.



During the EDA process we noticed that the data was highly imbalanced, with most of the flights taking off on time and only a handful of flights that were delayed and cancelled which made making sure our model classified flights as critically important to determining schedules for

transporting cargo. Figure 2 below, shows the distribution of flight statuses for the entire dataset.



Tools Used to Ingest Data

The primary tools for retrieving and preparing the data revolved around Amazon S3 buckets to store data and SageMaker notebooks to run python libraries and access AWS native services. Python libraries such as Pandas and NumPy were used to read in the CSV files directly from the S3 buckets, handle data manipulation and exploration and perform merged, until the point we load the processed data to be read by Amazon Athena to query directly from the processed_data bucket and bring it into the modeling phase. In future iterations we can leverage Amazon Athena and extend its application to the data cleaning phase of the project. In order to visualize our data and gain insights into the dataset Matplotlib and Seaborn were used

to develop insightful visualizations. The models were developed using SageMaker resources `sagemaker.sklearn.estimator` was used to develop the XGBoost classification model that we will use to predict whether or not the flight will be delayed or cancelled.

Measuring Impact

1. Prediction Accuracy of Flight Delays: The machine learning model should achieve at least 80% accuracy in predicting flights that are delayed/canceled. This can be measured by using the F1 score and precision between delayed and on-time flights.
2. Reduction in Operational Costs to Delays: Provide insights to reduce average cargo flight delay times by 2%. This will be measured by comparing the average delay duration before and after implementing predictive insights. Using Athena queries will help with tracking monthly delay trends and improvements.
3. Establishing Correlation between weather conditions and flight delays. The metric we will be using here is to identify the strongest weather-related contributor to flight delays at San Diego International Airport. We will perform a regression analysis to quantify the impact of weather conditions such as wind speed, temperature, etc. Also, a coefficient for each weather variance is developed to determine what influences delay probability.

Security Checklist, Privacy, and Other Risks

None of the data collected and stored for this project would be considered PHI, PII, or sensitive information, such as credit card numbers and user behavior. As data scientists, we are cognizant of bias making its way into our model. Since our model focuses on predicting flight delays specifically for San Diego International Airport, we want to avoid sampling, selection, temporal, omitted variable, and algorithmic biases. Addressing these is crucial to the

implementation and underlying trustworthiness of our model. We have addressed many of these biases by using National standardized governmental reports and data sources. By using the centralized data of every single flight departing from San Diego, we ensure sampling and selection bias are avoided through the inclusion of all data points and standardized time delays. We also aim to minimize omitted variable bias by including standardized weather conditions that were taken from the San Diego International Airport runway. By including this weather information, we aim to develop a deeper understanding of the weather phenomena that may impact delays and cancellations for San Diego International Airport.

S3 Buckets that our application will utilize to read and write data include a raw data bucket, which contains the raw weather and flight data spanning the two years. Additionally, we would like to have a separate bucket for the processed data that will be ready to feed into the model. Furthermore, once we are ready to develop our models, we would like to establish a model artifacts bucket that will contain the trained machine learning models and information on Hyperparameter Tuning and evaluation metrics of the model to quickly reference. Finally, we will have a bucket that will house the predictions and results predicting which flights will and will not be delayed.

No ethical concerns exist at this time with either the data being used or the scope of the project. We will make sure to update you if some unforeseen concern arises down the line.

Data Preparation

Firstly, the weather data needed to be extracted from its original file format. The METAR library is a library that is used to extract METAR values from a set format. We first had to isolate the REM column, which contained the rawest form of the METAR and SPECI data. Once we

isolated that column, we applied REGEX operations to break up the date time fields as well as the METAR data, after which we then applied the METAR library to extract station information, wind direction, wind speed, wind gusts, visibility, temperature, dewpoint, and altimeter setting. Once these columns had been extracted, we were free to drop the full metar code as the information had already been extracted. Missing wind gusts speeds were filled in with 0, indicative of no wind gusts at that time. The raw organized flight data was then merged with the organized weather data on a created single datetime variable. For instances where the flight data had no exact weather data match, the most recent weather observation was used before the flight took off. This is representative of the weather conditions at that time of takeoff. We handled missing values by filling DEP_TIME with 0 and using forward-fill for missing weather data.. Additionally, we ensured consistency in date/time formats between flight and weather data by removing seconds from the weather time field.

Feature Selection, Creation, and Transformation

We will use key features related to flight performance (FL_DATE, DEP_TIME, DEP_DEL15, CANCELLED), weather conditions (wind_direction, wind_speed_kt, temperature_c, dewpoint_c, visibility_statute_mi, altimeter_hg). Features like CANCELLATION_CODE will be excluded since CANCELLED already captures flight cancellations.

Our target variable column must be created using a combination of the binary categorical variable DEL_15 and CANCELLED. This target variable will be either a 1 for if the flight was delayed or cancelled and 0 if the flight left on time.

To fully get the data ready for modeling in the processed data set, we converted the categorical features found in the 'OP_UNIQUE_CARRIER' column along with the 'DEST' column using the "One-hot-encoding" technique to make these categorical features ready to be interpreted by the model.

Class Imbalance

As mentioned previously, our dataset was highly unbalanced and so needed to be balanced by undersampling the majority class. This is because we have a dataset with 185,600 observations, which allows us the ability to lose some data from our majority class. This will also allow us to establish a baseline model on which we can compare by also developing an oversampled model on which to compare and determine which balancing of data sets yields the best predictive results.

Train, Test, Validation Splits

Our dataset was split into thirds: a training dataset, a validation dataset, and a testing dataset. The split percentages will be 70/15/15, respectively. This maintains a sizable training portion of the data while also maintaining equal portions for the validation and testing sections of the dataset.

Model Training

This project uses a built-in SageMaker algorithm (XGBoost) rather than Jumpstart, a custom script, or a container. The decision was made after initial attempts to train a custom scikit-learn model encountered permissions issues and instance type restrictions. Switching to a built-in algorithm allowed the training job to be completed successfully within the constraints of the current AWS environment, reducing friction and enabling faster iteration.

Built-in Algorithm Approach

The final model uses XGBoost, a built-in SageMaker algorithm known for its speed, scalability, and high accuracy. While the initial approach used a Random Forest model built with scikit-learn, the team transitioned to XGBoost to take advantage of its optimized performance within the SageMaker environment. This shift supports the business goal of proactively identifying flight delays and cancellations by enabling faster training times and more efficient model deployment, which are critical for scaling predictive solutions in a production setting.

Data Training

During training, the train.csv and validation.csv datasets were passed in CSV format using the TrainingInput() function. The training instance type was set to ml.m5.large, based on available SageMaker quotas. While no custom hyperparameters were defined in this version, the selected parameters ensured compatibility with SageMaker's built-in XGBoost container and provided a stable environment to evaluate model performance.

Instance Size and Count

We used the ml.m5.large instance because it offers a balance of computing power and cost. It comes with 4 vCPUs and 16 GB of RAM, which is more than enough to handle our dataset of around 185,600 rows. It's also a more budget-friendly option compared to larger or GPU-based instances. We also considered scalability and resource efficiency, and overall, this instance size gives us a good mix of performance and affordability for what we're trying to do.

Model Evaluation

The model will be evaluated using recall as the primary metric, as it directly aligns with the business goal of proactively identifying flight delays and cancellations. Since the model is now trained using SageMaker's built-in XGBoost algorithm, recall remains a critical focus to

ensure that as many actual delays as possible are correctly identified. This minimizes operational disruptions and helps mitigate unplanned logistical costs. We chose to prioritize recall because missing a potential delay is far more costly than flagging a delay that doesn't happen. It is more significant for the business to catch true delays early.

Although precision is considered secondary, monitoring it ensures the model doesn't create unnecessary false positives that could overwhelm logistics teams. Additionally, the F1-score will be used as a balanced metric to summarize the model's overall performance, helping evaluate the trade-off between recall and precision in a single, interpretable value.

Overall, the model was successful in meeting the project's goals. It achieved a 95% recall, which supports planning and early delay detection. At a threshold of 0.25, the model reached a recall of 0.95 for the delayed/canceled class, which means it was able to catch most of the delays we were trying to identify. The model showed the connection between weather conditions and delays and supports the goal of helping logistics teams reduce average delay times with better planning.

Future Enhancements

Several key enhancements could bring this project closer to real-world application. Deploying the model with a SageMaker endpoint would provide fast, actionable insight that batch predictions just can't offer—making it a much better fit for day-to-day decision-making. It would allow airlines and logistics teams to respond immediately to high-risk flights, helping reduce delays and cancellations. Of course, this would depend on having the right IAM permissions and available resources, but it would be a strong upgrade.

Another major improvement would be automating the ML workflow with SageMaker Pipelines. Right now, much of the process—like uploading data, training the model, and triggering deployment—is manual. Automating it saves time, reduces human error, and ensures the model stays current with the latest data. In a business environment where conditions change constantly, having a consistent and repeatable pipeline would be an enhancement.

We could also improve the model's performance by integrating external data sources such as airport congestion and air traffic volume. These are real-world factors that heavily influence delays, and incorporating them would help the model predict disruptions more accurately. Given more time and resources, we could integrate external data via APIs, clean it, and add more features to enhance recall and reduce false positives—helping reduce disruptions and support business goals.

References

Amazon Web Services. (n.d.). *Deploy a scikit-learn model in Amazon SageMaker*. Retrieved from <https://docs.aws.amazon.com/sagemaker/latest/dg/sklearn.html>

Amazon Web Services. (n.d.). *SageMaker Pipelines: Build, automate, and manage ML workflows at scale*. Retrieved from <https://aws.amazon.com/sagemaker/pipelines/>

Bureau of Transportation Statistics. (n.d.). *On-Time: Reporting Carrier On-Time Performance*. U.S. Department of Transportation. Retrieved from https://www.transtats.bts.gov/Fields.asp?gnoyr_VQ=FGJ

National Centers for Environmental Information. (n.d.). *Integrated Surface Database (ISD)*.

National Oceanic and Atmospheric Administration (NOAA). Retrieved from

<https://www.ncei.noaa.gov/products/land-based-station/integrated-surface-database>

Appendix

Setup All Project Dependencies

****This code may take some time to run, please be patient ****

The following commands install the latest version of the packages.

In [1]: *# check python version and list the currently installed packages*

```
!python --version  
!pip list
```


Python 3.11.11

Package	Version
absl-py	2.1.0
accelerate	0.34.2
adagio	0.2.6
aiobotocore	2.20.0
aiohttp	3.9.5
aiohttp-cors	0.7.0
aioitertools	0.12.0
aiosignal	1.3.2
aiosqlite	0.19.0
alembic	1.15.1
altair	5.5.0
amazon-q-developer-jupyterlab-ext	3.4.7
amazon_sagemaker_jupyter_ai_q_developer	1.0.16
amazon_sagemaker_jupyter_scheduler	3.1.10
amazon-sagemaker-sql-editor	0.1.15
amazon-sagemaker-sql-execution	0.1.6
amazon-sagemaker-sql-magic	0.1.3
annotated-types	0.7.0
ansi2html	1.9.2
ansicolors	1.1.8
antlr4-python3-runtime	4.9.3
anyio	4.8.0
appdirs	1.4.4
archspec	0.2.5
argon2-cffi	23.1.0
argon2-cffi-bindings	21.2.0
arrow	1.3.0
asn1crypto	1.5.1
astroid	3.3.8
asttokens	3.0.0
astunparse	1.6.3
async-lru	2.0.4
async-timeout	4.0.3
attrs	23.2.0
autogluon	1.2
autogluon.common	1.2
autogluon.core	1.2
autogluon.features	1.2
autogluon.multimodal	1.2
autogluon.tabular	1.2
autogluon.timeseries	1.2
autopep8	2.0.4
autovizwidget	0.22.0

aws-embedded-metrics	3.2.0
aws-glue-sessions	1.0.8
babel	2.17.0
bcrypt	4.3.0
beautifulsoup4	4.13.3
binaryornot	0.4.4
bleach	6.2.0
blinker	1.9.0
blis	1.0.1
boltons	24.0.0
boto3	1.36.23
botocore	1.36.23
Brotli	1.1.0
cached-property	1.5.2
cachetools	5.5.2
catalogue	2.0.10
catboost	1.2.7
certifi	2025.1.31
cffi	1.17.1
chardet	5.2.0
charset-normalizer	3.4.1
click	8.1.8
cloudpathlib	0.20.0
cloudpickle	3.1.1
colorama	0.4.6
colorful	0.5.6
colorlog	6.9.0
comm	0.2.2
conda	24.11.3
conda-libmamba-solver	24.9.0
conda-package-handling	2.4.0
conda_package_streaming	0.11.0
confection	0.1.5
contextlib2	21.6.0
contourpy	1.3.1
cookiecutter	2.6.0
coreforecast	0.0.12
croniter	1.4.1
cryptography	44.0.2
cycler	0.12.1
cymem	2.0.11
cytoolz	1.0.1
dash	2.18.1
dask	2025.2.0
databricks-sdk	0.44.1
dataclasses	0.8

dataclasses-json	0.6.7
datasets	2.2.1
debugpy	1.8.13
decorator	5.2.1
deepmerge	2.0
defusedxml	0.7.1
Deprecated	1.2.18
dill	0.3.9
diskcache	5.6.3
distlib	0.3.9
distributed	2025.2.0
distro	1.9.0
dnspython	2.7.0
docker	7.1.0
docstring-to-markdown	0.15
email_validator	2.2.0
entrypoints	0.4
evaluate	0.4.1
exceptiongroup	1.2.2
executing	2.1.0
faiss	1.9.0
fastai	2.7.18
fastapi	0.115.11
fastapi-cli	0.0.7
fastcore	1.7.20
fastdownload	0.0.7
fastjsonschema	2.21.1
fastprogress	1.0.3
filelock	3.17.0
flake8	7.1.2
Flask	3.1.0
flatbuffers	25.2.10
fonttools	4.56.0
fqdn	1.5.1
frozendict	2.4.6
frozenset	1.5.0
fs	2.4.16
fsspec	2024.10.0
fugue	0.9.1
future	1.0.0
gast	0.6.0
gdown	5.2.0
git-remote-codecommit	1.16
gitdb	4.0.12
GitPython	3.1.44
gluonts	0.16.0

gmpy2	2.1.5
google-api-core	2.24.1
google-auth	2.38.0
google-pasta	0.2.0
googleapis-common-protos	1.69.0
graphene	3.4.3
graphql-core	3.2.6
graphql-relay	3.2.0
graphviz	0.20.3
greenlet	3.1.1
grpcio	1.62.2
gssapi	1.9.0
gunicorn	23.0.0
h11	0.14.0
h2	4.2.0
h5py	3.13.0
hdiupyterutils	0.22.0
hpack	4.1.0
httpcore	1.0.7
httptools	0.6.4
httpx	0.28.1
httpx-sse	0.4.0
huggingface_hub	0.29.1
hyperframe	6.1.0
hyperopt	0.2.7
idna	3.10
imagecodecs-lite	2019.12.3
imageio	2.37.0
importlib-metadata	6.10.0
importlib_resources	6.5.2
ipykernel	6.29.5
ipython	8.32.0
ipywidgets	8.1.5
isoduration	20.11.0
isort	6.0.1
itsdangerous	2.2.0
jedi	0.19.2
Jinja2	3.1.5
jmespath	1.0.1
joblib	1.4.2
json5	0.10.0
jsonpatch	1.33
jsonpath-ng	1.6.1
jsonpointer	3.0.0
jsonschema	4.23.0
jsonschema-specifications	2024.10.1

jupyter	1.1.1
jupyter-activity-monitor-extension	0.3.1
jupyter_ai	2.29.1
jupyter_ai_magics	2.29.1
jupyter_client	8.6.3
jupyter_collaboration	2.1.5
jupyter-console	6.6.3
jupyter_core	5.7.2
jupyter-dash	0.4.2
jupyter-events	0.12.0
jupyter-lsp	2.2.5
jupyter_scheduler	2.10.0
jupyter_server	2.15.0
jupyter_server_fileid	0.9.2
jupyter_server_mathjax	0.2.6
jupyter_server_proxy	4.4.0
jupyter_server_terminals	0.5.3
jupyter-ydoc	2.1.5
jupyterlab	4.3.5
jupyterlab_git	0.50.2
jupyterlab-lsp	5.0.3
jupyterlab_pygments	0.3.0
jupyterlab_server	2.27.3
jupyterlab_widgets	3.0.13
keras	3.8.0
kiwisolver	1.4.7
krb5	0.5.1
langchain	0.3.20
langchain-aws	0.2.10
langchain-community	0.3.19
langchain-core	0.3.41
langchain-text-splitters	0.3.6
langcodes	3.4.1
langsmith	0.2.11
language_data	1.3.0
lazy_loader	0.4
libmambapy	1.5.12
lightgbm	4.6.0
lightning	2.5.0.post0
lightning-utilities	0.13.1
linkify-it-py	2.0.3
llvmlite	0.44.0
locket	1.0.0
lxml	5.3.1
Mako	1.3.9
marisa-trie	1.2.1

Markdown	3.6
markdown-it-py	3.0.0
MarkupSafe	3.0.2
marshmallow	3.26.1
matplotlib	3.10.1
matplotlib-inline	0.1.7
mccabe	0.7.0
mdit-py-plugins	0.4.2
mdurl	0.1.2
memray	1.15.0
menuinst	2.2.0
mistune	3.1.2
ml-dtypes	0.4.0
mlflow	2.20.3
mlflow-skinny	2.20.3
mlforecast	0.13.4
mock	4.0.3
model-index	0.1.11
mpmath	1.3.0
msgpack	1.1.0
multidict	6.1.0
multiprocess	0.70.17
munkres	1.1.4
murmurhash	1.0.10
mypy_extensions	1.0.0
namex	0.0.8
narwhals	1.29.0
nbclient	0.10.2
nbconvert	7.16.6
nbdime	4.0.2
nbformat	5.10.4
nest_asyncio	1.6.0
networkx	3.4.2
nlpaug	1.1.11
nltk	3.9.1
nose	1.3.7
notebook	7.3.2
notebook_shim	0.2.4
numba	0.61.0
numpy	1.26.4
omegaconf	2.3.0
opencensus	0.11.3
opencensus-context	0.1.3
openmim	0.3.7
opentelemetry-api	1.30.0
opentelemetry-sdk	1.30.0

opentelemetry-semantic-conventions	0.51b0
opt_einsum	3.4.0
optree	0.14.1
optuna	4.2.1
ordered-set	4.1.0
orjson	3.10.15
overrides	7.7.0
packaging	24.2
pandas	2.2.3
pandocfilters	1.5.0
papermill	2.6.0
paramiko	3.5.1
parso	0.8.4
partd	1.4.2
pathos	0.3.3
patsy	1.0.1
pdf2image	1.17.0
pexpect	4.9.0
pickleshare	0.7.5
pillow	11.1.0
pip	24.3.1
pkgutil_resolve_name	1.3.10
platformdirs	4.3.6
plotly	6.0.0
pluggy	1.5.0
ply	3.11
pox	0.3.5
ppft	1.7.6.9
preshed	3.0.9
prometheus_client	0.21.1
prometheus_flask_exporter	0.23.1
prompt_toolkit	3.0.50
propcache	0.2.1
proto-plus	1.26.0
protobuf	4.25.3
psutil	5.9.8
ptyprocess	0.7.0
pure_eval	0.2.3
pure-sasl	0.6.2
py4j	0.10.9.9
pyarrow	17.0.0
pyasn1	0.6.1
pyasn1_modules	0.4.1
PyAthena	3.12.2
pycodestyle	2.12.1
pycosat	0.6.6

pycparser	2.22
pycrdt	0.12.8
pycrdt-websocket	0.15.4
pydantic	2.10.6
pydantic_core	2.27.2
pydantic-settings	2.8.1
pydocstyle	6.3.0
pyflakes	3.2.0
Pygments	2.19.1
PyHive	0.7.0
PyJWT	2.10.1
pylint	3.3.4
PyNaCl	1.5.0
pyOpenSSL	24.3.0
pyparsing	3.2.1
PySocks	1.7.1
pyspnego	0.11.2
pytesseract	0.3.10
python-dateutil	2.9.0.post0
python-dotenv	1.0.1
python-json-logger	2.0.7
python-lsp-jsonrpc	1.1.2
python-lsp-server	1.12.2
python-multipart	0.0.20
python-slugify	8.0.4
pytoolconfig	1.2.5
pytorch-lightning	2.5.0.post0
pytorch-metric-learning	2.3.0
pytz	2024.1
pyu2f	0.1.5
PyWavelets	1.8.0
PyYAML	6.0.2
pyzmq	26.2.1
querystring_parser	1.2.4
ray	2.37.0
redshift_connector	2.1.5
referencing	0.36.2
regex	2024.11.6
requests	2.32.3
requests-kerberos	0.15.0
requests-toolbelt	1.0.0
responses	0.18.0
retrying	1.3.4
rfc3339_validator	0.1.4
rfc3986-validator	0.1.1
rich	13.9.4

rich-toolkit	0.11.3
rope	1.13.0
rpds-py	0.23.1
rsa	4.9
ruamel.yaml	0.18.10
ruamel.yaml.clib	0.2.8
s3fs	2024.10.0
s3transfer	0.11.3
safetensors	0.5.3
sagemaker	2.240.0
sagemaker-core	1.0.25
sagemaker-headless-execution-driver	0.0.13
sagemaker-jupyterlab-emr-extension	0.3.7
sagemaker-jupyterlab-extension	0.4.0
sagemaker-jupyterlab-extension-common	0.1.36
sagemaker-kernel-wrapper	0.0.5
sagemaker-mlflow	0.1.0
sagemaker-studio-analytics-extension	0.1.5
sagemaker-studio-sparkmagic-lib	0.1.4
schema	0.7.7
scikit-image	0.20.0
scikit-learn	1.5.2
scipy	1.15.2
scramp	1.4.4
seaborn	0.13.2
Send2Trash	1.8.3
sentencepiece	0.1.99
sequeval	1.2.2
setproctitle	1.3.5
setuptools	75.8.2
shellingham	1.5.4
simpervisor	1.0.0
six	1.17.0
smart_open	7.1.0
smdebug-rulesconfig	1.0.1
smmap	5.0.2
sniffio	1.3.1
snowballstemmer	2.2.0
snowflake-connector-python	3.13.2
sortedcontainers	2.4.0
soupsieve	2.5
spacy	3.8.2
spacy-legacy	3.0.12
spacy-loggers	1.0.5
sparkmagic	0.21.0
SQLAlchemy	2.0.38

sqlite-anyio	0.2.3
sqlparse	0.5.3
srsly	2.5.1
stack_data	0.6.3
starlette	0.46.0
statsforecast	1.7.8
statsmodels	0.14.4
supervisor	4.2.5
sympy	1.13.3
tabulate	0.9.0
tblib	3.0.0
tenacity	9.0.0
tensorboard	2.17.1
tensorboard_data_server	0.7.0
tensorboardX	2.6.2.2
tensorflow	2.17.0
tensorflow_estimator	2.15.0
termcolor	2.5.0
terminado	0.18.1
text-unidecode	1.3
textual	2.1.2
tf_keras	2.17.0
thinc	8.3.2
threadpoolctl	3.5.0
thrift	0.20.0
thrift_sasl	0.4.3
tifffile	2020.6.3
timm	1.0.3
tinycss2	1.4.0
tokenizers	0.21.0
tomli	2.2.1
tomlkit	0.13.2
toolz	0.12.1
torch	2.4.1.post100
torchmetrics	1.2.1
torchvision	0.19.1
tornado	6.4.2
tqdm	4.67.1
traitlets	5.14.3
transformers	4.49.0
triad	0.9.8
truststore	0.10.1
typer	0.15.2
typer-slim	0.15.2
types-python-dateutil	2.9.0.20241206
typing_extensions	4.12.2

typing_inspect	0.9.0
typing_utils	0.1.0
tzdata	2025.1
uc-micro-py	1.0.3
ujson	5.10.0
unicodedata2	16.0.0
uri-template	1.3.0
urllib3	2.3.0
utilsforecast	0.2.3
uvicorn	0.34.0
uvloop	0.21.0
virtualenv	20.29.2
wasabi	1.1.3
watchfiles	0.24.0
wcwidth	0.2.13
weasel	0.4.1
webcolors	24.11.1
webencodings	0.5.1
websocket-client	1.8.0
websockets	15.0
Werkzeug	3.1.3
whatthepatch	1.0.7
wheel	0.45.1
widetsnbextension	4.0.13
window_ops	0.0.15
wrapt	1.17.2
xgboost	2.1.4
xxhash	3.5.0
y-py	0.6.2
yapf	0.43.0
yaml	1.18.3
ypy-websocket	0.12.4
zict	3.0.0
zipp	3.21.0
zstandard	0.19.0

```
In [2]: # Upgrade pip and wrapt to the latest versions
!pip install --disable-pip-version-check -q pip --upgrade > /dev/null
!pip install --disable-pip-version-check -q wrapt --upgrade > /dev/null
```

AWS CLI and AWS Python SDK (boto3)

This installs the latest versions of AWS CLI, boto3 and botocore

```
In [3]: !pip install --disable-pip-version-check -q awscli boto3 botocore
```

```
ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the
source of the following dependency conflicts.
autogluon-multimodal 1.2 requires nvidia-ml-py3==7.352.0, which is not installed.
aiobotocore 2.20.0 requires botocore<1.36.24,>=1.36.20, but you have botocore 1.37.34 which is incompatible.
amazon-sagemaker-sql-magic 0.1.3 requires sqlparse==0.5.0, but you have sqlparse 0.5.3 which is incompatible.
autogluon-multimodal 1.2 requires jsonschema<4.22,>=4.18, but you have jsonschema 4.23.0 which is incompatible.
autogluon-multimodal 1.2 requires nltk<3.9,>=3.4.5, but you have nltk 3.9.1 which is incompatible.
autogluon-multimodal 1.2 requires omegaconf<2.3.0,>=2.1.1, but you have omegaconf 2.3.0 which is incompatible.
```

SageMaker Services

This section installs the latest version of sagemaker, smdebug, and sagemaker-experiments

```
In [15]: !pip install --disable-pip-version-check -q sagemaker smdebug sagemaker-experiments
```

PyTorch

This section installs PyTorch using conda for compatibility

```
In [16]: !conda install -y pytorch -c pytorch
```

Channels:

- pytorch
- conda-forge

Platform: linux-64

doneecting package metadata (repodata.json): -

^Clving environment: /

TensorFlow

Installs the latest version of TensorFlow

```
In [17]: !pip install --disable-pip-version-check -q tensorflow
```

^C

ERROR: Operation cancelled by user

PyAthena

Installs the latest version of PyAthena

```
In [18]: !pip install --disable-pip-version-check -q PyAthena
```

AWS Data Wrangler

Installs the latest version of AWS Data Wrangler

```
In [19]: !pip install --disable-pip-version-check -q awswrangler
```

Matplotlib, Seaborn, Pandas

This installs the latest versions of Matplotlib, Seaborn, Pandas

```
In [20]: !pip install --disable-pip-version-check -q matplotlib
!pip install --disable-pip-version-check -q seaborn
!pip install --disable-pip-version-check -q pandas
```

METAR Library

This will install the metar library which is used to read in raw formatted metar data in order to convert the values into columns

```
In [21]: !pip install --disable-pip-version-check -q metar
```

Windrose visualization library

This will install a library that aids in the development of wind rose visualizations showing where wind direction and speed are common.

```
In [1]: !pip install --disable-pip-version-check -q windrose
```

Checking Final versions of python and packages

```
In [23]: !python --version
!pip list
```

Python 3.11.11

Package	Version
absl-py	2.1.0
accelerate	0.34.2
adagio	0.2.6
aiobotocore	2.20.0
aiohttp	3.9.5
aiohttp-cors	0.7.0
aioitertools	0.12.0
aiosignal	1.3.2
aiosqlite	0.19.0
alembic	1.15.1
altair	5.5.0
amazon-q-developer-jupyterlab-ext	3.4.7
amazon_sagemaker_jupyter_ai_q_developer	1.0.16
amazon_sagemaker_jupyter_scheduler	3.1.10
amazon-sagemaker-sql-editor	0.1.15
amazon-sagemaker-sql-execution	0.1.6
amazon-sagemaker-sql-magic	0.1.3
annotated-types	0.7.0
ansi2html	1.9.2
ansicolors	1.1.8
antlr4-python3-runtime	4.9.3
anyio	4.8.0
appdirs	1.4.4
archspec	0.2.5
argon2-cffi	23.1.0
argon2-cffi-bindings	21.2.0
arrow	1.3.0
asn1crypto	1.5.1
astroid	3.3.8
asttokens	3.0.0
astunparse	1.6.3
async-lru	2.0.4
async-timeout	4.0.3
attrs	23.2.0
autogluon	1.2
autogluon.common	1.2
autogluon.core	1.2
autogluon.features	1.2
autogluon.multimodal	1.2
autogluon.tabular	1.2
autogluon.timeseries	1.2
autopep8	2.0.4
autovizwidget	0.22.0

aws-embedded-metrics	3.2.0
aws-glue-sessions	1.0.8
awscli	1.38.34
awswrangler	3.11.0
babel	2.17.0
bcrypt	4.3.0
beautifulsoup4	4.13.3
binaryornot	0.4.4
bleach	6.2.0
blinker	1.9.0
blis	1.0.1
boltons	24.0.0
boto3	1.37.34
botocore	1.37.34
Brotli	1.1.0
cached-property	1.5.2
cachetools	5.5.2
catalogue	2.0.10
catboost	1.2.7
certifi	2025.1.31
cffi	1.17.1
charset	5.2.0
charset-normalizer	3.4.1
click	8.1.8
cloudpathlib	0.20.0
cloudpickle	3.1.1
colorama	0.4.6
colorful	0.5.6
colorlog	6.9.0
comm	0.2.2
conda	24.11.3
conda-libmamba-solver	24.9.0
conda-package-handling	2.4.0
conda_package_streaming	0.11.0
confection	0.1.5
contextlib2	21.6.0
contourpy	1.3.1
cookiecutter	2.6.0
coreforecast	0.0.12
croniter	1.4.1
cryptography	44.0.2
cycler	0.12.1
cymem	2.0.11
cytoolz	1.0.1
dash	2.18.1
dask	2025.2.0

databricks-sdk	0.44.1
dataclasses	0.8
dataclasses-json	0.6.7
datasets	2.2.1
debugpy	1.8.13
decorator	5.2.1
deepmerge	2.0
defusedxml	0.7.1
Deprecated	1.2.18
dill	0.3.9
diskcache	5.6.3
distlib	0.3.9
distributed	2025.2.0
distro	1.9.0
dnspython	2.7.0
docker	7.1.0
docstring-to-markdown	0.15
docutils	0.19
email_validator	2.2.0
entrypoints	0.4
evaluate	0.4.1
exceptiongroup	1.2.2
executing	2.1.0
faiss	1.9.0
fastai	2.7.18
fastapi	0.115.11
fastapi-cli	0.0.7
fastcore	1.7.20
fastdownload	0.0.7
fastjsonschema	2.21.1
fastprogress	1.0.3
filelock	3.17.0
flake8	7.1.2
Flask	3.1.0
flatbuffers	25.2.10
fonttools	4.56.0
fqdn	1.5.1
frozendict	2.4.6
frozenset	1.5.0
fs	2.4.16
fsspec	2024.10.0
fugue	0.9.1
future	1.0.0
gast	0.6.0
gdown	5.2.0
git-remote-codecommit	1.16

gitdb	4.0.12
GitPython	3.1.44
gluons	0.16.0
gmpy2	2.1.5
google-api-core	2.24.1
google-auth	2.38.0
google-pasta	0.2.0
googleapis-common-protos	1.69.0
graphene	3.4.3
graphql-core	3.2.6
graphql-relay	3.2.0
graphviz	0.20.3
greenlet	3.1.1
grpcio	1.62.2
gssapi	1.9.0
gunicorn	23.0.0
h11	0.14.0
h2	4.2.0
h5py	3.13.0
hdiupyterutils	0.22.0
hpack	4.1.0
httpcore	1.0.7
httptools	0.6.4
httpx	0.28.1
httpx-sse	0.4.0
huggingface_hub	0.29.1
hyperframe	6.1.0
hyperopt	0.2.7
idna	3.10
imagecodecs-lite	2019.12.3
imageio	2.37.0
importlib-metadata	6.10.0
importlib_resources	6.5.2
ipykernel	6.29.5
ipython	8.32.0
ipywidgets	8.1.5
isoduration	20.11.0
isort	6.0.1
itsdangerous	2.2.0
jedi	0.19.2
Jinja2	3.1.5
jmespath	1.0.1
joblib	1.4.2
json5	0.10.0
jsonpatch	1.33
jsonpath-ng	1.6.1

jsonpointer	3.0.0
jsonschema	4.23.0
jsonschema-specifications	2024.10.1
jupyter	1.1.1
jupyter-activity-monitor-extension	0.3.1
jupyter_ai	2.29.1
jupyter_ai_magics	2.29.1
jupyter_client	8.6.3
jupyter_collaboration	2.1.5
jupyter-console	6.6.3
jupyter_core	5.7.2
jupyter-dash	0.4.2
jupyter-events	0.12.0
jupyter-lsp	2.2.5
jupyter_scheduler	2.10.0
jupyter_server	2.15.0
jupyter_server_fileid	0.9.2
jupyter_server_mathjax	0.2.6
jupyter_server_proxy	4.4.0
jupyter_server_terminals	0.5.3
jupyter-ydoc	2.1.5
jupyterlab	4.3.5
jupyterlab_git	0.50.2
jupyterlab-lsp	5.0.3
jupyterlab_pygments	0.3.0
jupyterlab_server	2.27.3
jupyterlab_widgets	3.0.13
keras	3.8.0
kiwisolver	1.4.7
krb5	0.5.1
langchain	0.3.20
langchain-aws	0.2.10
langchain-community	0.3.19
langchain-core	0.3.41
langchain-text-splitters	0.3.6
langcodes	3.4.1
langsmith	0.2.11
language_data	1.3.0
lazy_loader	0.4
libmambapy	1.5.12
lightgbm	4.6.0
lightning	2.5.0.post0
lightning-utilities	0.13.1
linkify-it-py	2.0.3
llvmlite	0.44.0
locket	1.0.0

lxml	5.3.1
Mako	1.3.9
marisa-trie	1.2.1
Markdown	3.6
markdown-it-py	3.0.0
MarkupSafe	3.0.2
marshmallow	3.26.1
matplotlib	3.10.1
matplotlib-inline	0.1.7
mccabe	0.7.0
mdit-py-plugins	0.4.2
mdurl	0.1.2
memray	1.15.0
menuinst	2.2.0
metar	1.11.0
mistune	3.1.2
ml-dtypes	0.4.0
mlflow	2.20.3
mlflow-skinny	2.20.3
mlforecast	0.13.4
mock	4.0.3
model-index	0.1.11
mpmath	1.3.0
msgpack	1.1.0
multidict	6.1.0
multiprocess	0.70.17
munkres	1.1.4
murmurhash	1.0.10
mypy_extensions	1.0.0
namex	0.0.8
narwhals	1.29.0
nbclient	0.10.2
nbconvert	7.16.6
nbdime	4.0.2
nbformat	5.10.4
nest_asyncio	1.6.0
networkx	3.4.2
nlpaug	1.1.11
nltk	3.9.1
nose	1.3.7
notebook	7.3.2
notebook_shim	0.2.4
numba	0.61.0
numpy	1.26.4
omegaconf	2.3.0
opencensus	0.11.3

opencensus-context	0.1.3
openmim	0.3.7
opentelemetry-api	1.30.0
opentelemetry-sdk	1.30.0
opentelemetry-semantic-conventions	0.51b0
opt_einsum	3.4.0
optree	0.14.1
optuna	4.2.1
ordered-set	4.1.0
orjson	3.10.15
overrides	7.7.0
packaging	24.2
pandas	2.2.3
pandocfilters	1.5.0
papermill	2.6.0
paramiko	3.5.1
parso	0.8.4
partd	1.4.2
pathos	0.3.3
patsy	1.0.1
pdf2image	1.17.0
pexpect	4.9.0
pickleshare	0.7.5
pillow	11.1.0
pip	25.0.1
pkgutil_resolve_name	1.3.10
platformdirs	4.3.6
plotly	6.0.0
pluggy	1.5.0
ply	3.11
pox	0.3.5
ppft	1.7.6.9
prashed	3.0.9
prometheus_client	0.21.1
prometheus_flask_exporter	0.23.1
prompt_toolkit	3.0.50
propcache	0.2.1
proto-plus	1.26.0
protobuf	3.20.3
psutil	5.9.8
ptyprocess	0.7.0
pure_eval	0.2.3
pure-sasl	0.6.2
py4j	0.10.9.9
pyarrow	17.0.0
pyasn1	0.6.1

pyasn1_modules	0.4.1
PyAthena	3.12.2
pycodestyle	2.12.1
pycosat	0.6.6
pycparser	2.22
pycrdt	0.12.8
pycrdt-websocket	0.15.4
pydantic	2.10.6
pydantic_core	2.27.2
pydantic-settings	2.8.1
pydocstyle	6.3.0
pyflakes	3.2.0
Pygments	2.19.1
PyHive	0.7.0
pyinstrument	3.4.2
pyinstrument_cext	0.2.4
PyJWT	2.10.1
pylint	3.3.4
PyNaCl	1.5.0
pyOpenSSL	24.3.0
pyparsing	3.2.1
PySocks	1.7.1
pyspnego	0.11.2
pytesseract	0.3.10
python-dateutil	2.9.0.post0
python-dotenv	1.0.1
python-json-logger	2.0.7
python-lsp-jsonrpc	1.1.2
python-lsp-server	1.12.2
python-multipart	0.0.20
python-slugify	8.0.4
pytoolconfig	1.2.5
pytorch-lightning	2.5.0.post0
pytorch-metric-learning	2.3.0
pytz	2024.1
pyu2f	0.1.5
PyWavelets	1.8.0
PyYAML	6.0.2
pyzmq	26.2.1
querystring_parser	1.2.4
ray	2.37.0
redshift_connector	2.1.5
referencing	0.36.2
regex	2024.11.6
requests	2.32.3
requests-kerberos	0.15.0

requests-toolbelt	1.0.0
responses	0.18.0
retrying	1.3.4
rfc3339_validator	0.1.4
rfc3986-validator	0.1.1
rich	13.9.4
rich-toolkit	0.11.3
rope	1.13.0
rpds-py	0.23.1
rsa	4.7.2
ruamel.yaml	0.18.10
ruamel.yaml.clib	0.2.8
s3fs	2024.10.0
s3transfer	0.11.3
safetensors	0.5.3
sagemaker	2.240.0
sagemaker-core	1.0.25
sagemaker-experiments	0.1.45
sagemaker-headless-execution-driver	0.0.13
sagemaker-jupyterlab-emr-extension	0.3.7
sagemaker-jupyterlab-extension	0.4.0
sagemaker-jupyterlab-extension-common	0.1.36
sagemaker-kernel-wrapper	0.0.5
sagemaker-mlflow	0.1.0
sagemaker-studio-analytics-extension	0.1.5
sagemaker-studio-sparkmagic-lib	0.1.4
schema	0.7.7
scikit-image	0.20.0
scikit-learn	1.5.2
scipy	1.15.2
scramp	1.4.4
seaborn	0.13.2
Send2Trash	1.8.3
sentencepiece	0.1.99
sequeval	1.2.2
setproctitle	1.3.5
setuptools	75.8.2
shellingham	1.5.4
simpervisor	1.0.0
six	1.17.0
smart_open	7.1.0
smdebug	1.0.34
smdebug-rulesconfig	1.0.1
smmmap	5.0.2
sniffio	1.3.1
snowballstemmer	2.2.0

snowflake-connector-python	3.13.2
sortedcontainers	2.4.0
soupsieve	2.5
spacy	3.8.2
spacy-legacy	3.0.12
spacy-loggers	1.0.5
sparkmagic	0.21.0
SQLAlchemy	2.0.38
sqlite-anyio	0.2.3
sqlparse	0.5.3
srsly	2.5.1
stack_data	0.6.3
starlette	0.46.0
statsforecast	1.7.8
statsmodels	0.14.4
supervisor	4.2.5
sympy	1.13.3
tabulate	0.9.0
tblib	3.0.0
tenacity	9.0.0
tensorboard	2.17.1
tensorboard_data_server	0.7.0
tensorboardX	2.6.2.2
tensorflow	2.17.0
tensorflow_estimator	2.15.0
termcolor	2.5.0
terminado	0.18.1
text-unidecode	1.3
textual	2.1.2
tf_keras	2.17.0
thinc	8.3.2
threadpoolctl	3.5.0
thrift	0.20.0
thrift_sasl	0.4.3
tifffile	2020.6.3
timm	1.0.3
tinycss2	1.4.0
tokenizers	0.21.0
tomli	2.2.1
tomlkit	0.13.2
toolz	0.12.1
torch	2.4.1.post100
torchmetrics	1.2.1
torchvision	0.19.1
tornado	6.4.2
tqdm	4.67.1

traitlets	5.14.3
transformers	4.49.0
triad	0.9.8
truststore	0.10.1
typer	0.15.2
typer-slim	0.15.2
types-python-dateutil	2.9.0.20241206
typing_extensions	4.12.2
typing_inspect	0.9.0
typing_utils	0.1.0
tzdata	2025.1
uc-micro-py	1.0.3
ujson	5.10.0
unicodedata2	16.0.0
uri-template	1.3.0
urllib3	2.3.0
utilsforecast	0.2.3
uvicorn	0.34.0
uvloop	0.21.0
virtualenv	20.29.2
wasabi	1.1.3
watchfiles	0.24.0
wcwidth	0.2.13
weasel	0.4.1
webcolors	24.11.1
webencodings	0.5.1
websocket-client	1.8.0
websockets	15.0
Werkzeug	3.1.3
whatthepatch	1.0.7
wheel	0.45.1
widgetsnbextension	4.0.13
window_ops	0.0.15
windrose	1.9.2
wrapt	1.17.2
xgboost	2.1.4
xxhash	3.5.0
y-py	0.6.2
yapf	0.43.0
yaml	1.18.3
ypy-websocket	0.12.4
zict	3.0.0
zipp	3.21.0
zstandard	0.19.0


```
In [24]: # Mark the dependency setup as complete
setup_dependencies_passed = True
%store setup_dependencies_passed
%store
```

Stored 'setup_dependencies_passed' (bool)

Stored variables and their in-db values:

```
setup_dependencies_passed      -> True
setup_iam_roles_passed         -> True
setup_instance_check_passed    -> True
setup_s3_bucket_passed         -> True
```

Release the resources

Following cells will shut down your kernel in order to release resources

```
In [2]: %%html

abs<p><b>Shutting down your kernel for this notebook to release resources.</b></p>
<button class="sm-command-button" data-commandlinker-command="kernelmenu:shutdown" style="display:none;">Shutdown Kernel</button>

<script>
try {
  els = document.getElementsByClassName("sm-command-button");
  els[0].click();
}
catch(err) {
  // NoOp
}
</script>
```

abs

Shutting down your kernel for this notebook to release resources.

In []:

Check Environment

In [1]:

```
import boto3

region = boto3.Session().region_name
session = boto3.Session()

ec2 = boto3.Session().client(service_name="ec2", region_name=region)
sm = boto3.Session().client(service_name="sagemaker", region_name=region)
```

In [2]:

```
import json

notebook_instance_name = None

try:
    with open("/opt/ml/metadata/resource-metadata.json") as notebook_info:
        data = json.load(notebook_info)
        domain_id = data["DomainId"]
        resource_arn = data["ResourceArn"]
        region = resource_arn.split(":")[3]
        name = data["ResourceName"]
        print("DomainId: {}".format(domain_id))
        print("Name: {}".format(name))
except:
    print("+++++")
    print("[ERROR]: COULD NOT RETRIEVE THE METADATA.")
    print("+++++")
```

DomainId: d-dlu97fb10zsw

Name: default

In [3]:

```
describe_domain_response = sm.describe_domain(DomainId=domain_id)
print(describe_domain_response["Status"])
```

InService

In [4]:

```
try:
    get_status_response = sm.get_sagemaker_servicecatalog_portfolio_status()
    print(get_status_response["Status"])
except:
    pass
```

Enabled

```
In [5]: if (
        describe_domain_response["Status"] == "InService"
        and get_status_response["Status"] == "Enabled"
        and "default" in name
    ):
        setup_instance_check_passed = True
        print("[OK] Checks passed! Great Job!! Please Continue.")
    else:
        setup_instance_check_passed = False
        print("+++++")
        print("[ERROR]: WE HAVE IDENTIFIED A MISCONFIGURATION.")
        print(describe_domain_response["Status"])
        print(get_status_response["Status"])
        print(name)
        print("+++++")
```

[OK] Checks passed! Great Job!! Please Continue.

```
In [6]: print(setup_instance_check_passed)
```

True

```
In [7]: %store setup_instance_check_passed
```

Stored 'setup_instance_check_passed' (bool)

```
In [8]: %store
```

Stored variables and their in-db values:

setup_dependencies_passed	-> True
setup_iam_roles_passed	-> True
setup_instance_check_passed	-> True
setup_s3_bucket_passed	-> True

```
In [1]: %%html
```

```
<p><b>Shutting down your kernel for this notebook to release resources.</b></p>
<button class="sm-command-button" data-commandlinker-command="kernelmenu:shutdown" style="display:none;">Shutdown Kernel</button>

<script>
try {
    els = document.getElementsByClassName("sm-command-button");
    els[0].click();
}
catch(err) {
```

```
// NoOp  
}  
</script>
```

Shutting down your kernel for this notebook to release resources.

In []:

Create S3 buckets for stages of the project

In order to establish these buckets in your own sagemaker environment, simply change the initials at the end of each bucket to your unique initials.

In [1]:

```
# import libraries
import boto3
import sagemaker

from botocore.client import ClientError
```

```
/opt/conda/lib/python3.11/site-packages/pydantic/_internal/_fields.py:192: UserWarning: Field name "json" in "MonitoringDatasetFormat" shadows an attribute in parent "Base"
  warnings.warn(
```

```
sagemaker.config INFO - Not applying SDK defaults from location: /etc/xdg/sagemaker/config.yaml
```

```
sagemaker.config INFO - Not applying SDK defaults from location: /home/sagemaker-user/.config/sagemaker/config.yaml
```

Establish session and client information

In [2]:

```
session = boto3.session.Session()
region = session.region_name
s3 = boto3.Session().client(service_name="s3", region_name=region)
```

Create Raw data S3 bucket

In [3]:

```
#set env variable
%env RAW_BUCKET=group9-ml-proj-raw-data-bucket-to

raw_bucket = "group9-ml-proj-raw-data-bucket-to"

try:
    if region == "us-east-1":
        s3.create_bucket(Bucket=raw_bucket)
    else:
        s3.create_bucket(Bucket=raw_bucket, CreateBucketConfiguration={"LocationConstraint": region})
    print("Raw Data Bucket created:", raw_bucket)
except Exception as e:
    print("Raw Data Bucket may already exist or an error occurred:", e)
```

```
env: RAW_BUCKET=group9-m1-proj-raw-data-bucket-grw
Raw Data Bucket created: group9-m1-proj-raw-data-bucket-grw
```

```
In [4]: %%bash
aws s3 ls s3://$RAW_BUCKET/
```

Create Processed data bucket

```
In [5]: #set env variable
%env PROCESSED_BUCKET=group9-m1-proj-processed-data-bucket-

processed_bucket = "group9-m1-proj-processed-data-bucket-grw"

try:
    if region == "us-east-1":
        s3.create_bucket(Bucket=processed_bucket)
    else:
        s3.create_bucket(Bucket=processed_bucket, CreateBucketConfiguration={"LocationConstraint": region})
    print("Processed Data Bucket created:", processed_bucket)
except Exception as e:
    print("Processed Data Bucket may already exist or an error occurred:", e)
```

```
env: PROCESSED_BUCKET=group9-m1-proj-processed-data-bucket-grw
Processed Data Bucket created: group9-m1-proj-processed-data-bucket-grw
```

```
In [6]: %%bash
aws s3 ls s3://$PROCESSED_BUCKET/
```

Athena Query results bucket

```
In [ ]: #set env variable
%env QUERY_BUCKET=group9-m1-proj-athena-query-bucket-grw

query_bucket = "group9-m1-proj-athena-query-bucket-grw"

try:
    if region == "us-east-1":
        s3.create_bucket(Bucket=query_bucket)
    else:
        s3.create_bucket(Bucket=query_bucket, CreateBucketConfiguration={"LocationConstraint": region})
    print("Athena Query Bucket created:", query_bucket)
```

```
except Exception as e:
    print("Athena Query Bucket may already exist or an error occurred:", e)
```

```
In [ ]: %%bash
aws s3 ls s3://$QUERY_BUCKET/
```

Create Training Artifacts bucket

```
In [7]: #set env variable
%env TNG_ART_BUCKET=group9-ml-proj-training-artifacts-bucket-grw

tng_art_bucket = "group9-ml-proj-training-artifacts-bucket-grw"

try:
    if region == "us-east-1":
        s3.create_bucket(Bucket=tng_art_bucket)
    else:
        s3.create_bucket(Bucket=tng_art_bucket, CreateBucketConfiguration={"LocationConstraint": region})
    print("Training Artifacts Bucket created:", tng_art_bucket)
except Exception as e:
    print("Training Artifacts Bucket may already exist or an error occurred:", e)
```

env: TNG_ART_BUCKET=group9-ml-proj-training-artifacts-bucket-grw
Training Artifacts Bucket created: group9-ml-proj-training-artifacts-bucket-grw

```
In [8]: %%bash
aws s3 ls s3://$TNG_ART_BUCKET/
```

Create Model Artifacts Bucket

```
In [9]: #set env variable
%env MONITORING_BUCKET=group9-ml-proj-model-artifacts-bucket-grw

model_art_bucket = "group9-ml-proj-model-artifacts-bucket-grw"

try:
    if region == "us-east-1":
        s3.create_bucket(Bucket=model_art_bucket)
    else:
        s3.create_bucket(Bucket=model_art_bucket, CreateBucketConfiguration={"LocationConstraint": region})
    print("Model Artifacts Bucket created:", model_art_bucket)
```

```
except Exception as e:
    print("Model Artifacts Bucket may already exist or an error occurred:", e)
```

env: MONITORING_BUCKET=group9-ml-proj-model-artifacts-bucket-grw
Model Artifacts Bucket created: group9-ml-proj-model-artifacts-bucket-grw

```
In [10]: %%bash
aws s3 ls s3://$MODEL_ART_BUCKET/
```

```
2025-03-24 17:48:48 group9-ml-proj-model-artifacts-bucket-grw
2025-03-24 17:48:46 group9-ml-proj-processed-data-bucket-grw
data-bucket-grw8:45 group9-ml-proj-raw-
2025-03-24 17:48:47 group9-ml-proj-training-artifacts-bucket-grw
sagemaker-studio-7sb8gxpljq9
2025-02-28 21:29:38 sagemaker-studio-chxawrkbs1q
agemaker-studio-o4i4ukj5y3b
2025-02-28 21:29:40 sagemaker-us-east-1-804823187130
```

Create Monitoring Bucket

```
In [11]: #set env variable
%env MONITORING_BUCKET=group9-ml-proj-monitoring-bucket-grw

monitoring_bucket = "group9-ml-proj-monitoring-bucket-grw"

try:
    if region == "us-east-1":
        s3.create_bucket(Bucket=monitoring_bucket)
    else:
        s3.create_bucket(Bucket=monitoring_bucket, CreateBucketConfiguration={"LocationConstraint": region})
    print("Inference/Monitoring Bucket created:", monitoring_bucket)
except Exception as e:
    print("Inference/Monitoring Bucket may already exist or an error occurred:", e)
```

env: MONITORING_BUCKET=group9-ml-proj-monitoring-bucket-grw
Inference/Monitoring Bucket created: group9-ml-proj-monitoring-bucket-grw

```
In [12]: %%bash
aws s3 ls s3://$MONITORING_BUCKET/
```

```
In [1]: %%html
<p><b>Shutting down your kernel for this notebook to release resources.</b></p>
<button class="sm-command-button" data-commandlinker-command="kernelmenu:shutdown" style="display:none;">Shutdown Kernel</button>
```



```
<script>
try {
  els = document.getElementsByClassName("sm-command-button");
  els[0].click();
}
catch(err) {
  // NoOp
}
</script>
```

Shutting down your kernel for this notebook to release resources.

In []:

Athena Database and Table Creation

This notebook will establish the AWS Athena Database that will be responsible for developing the raw_data and processed_data tables which will store the data.

```
In [1]: # import libraries
import boto3
import time
```

```
In [2]: # run athena query function
def run_athena_query(query, database, output_location):
    athena_client = boto3.client('athena')
    # query execution.
    response = athena_client.start_query_execution(
        QueryString=query,
        QueryExecutionContext={'Database': database},
        ResultConfiguration={'OutputLocation': output_location}
    )
    query_execution_id = response['QueryExecutionId']
    print(f"Query submitted. Execution ID: {query_execution_id}")
    # poll query until it completes.
    state = 'RUNNING'
    while state in ['RUNNING', 'QUEUED']:
        response = athena_client.get_query_execution(QueryExecutionId=query_execution_id)
        state = response['QueryExecution']['Status']['State']
        print(f"Current query state: {state}")
        if state in ['RUNNING', 'QUEUED']:
            time.sleep(2)

    if state == 'FAILED':
        reason = response['QueryExecution']['Status'].get('StateChangeReason', 'Unknown error')
        raise Exception(f"Query {query_execution_id} failed: {reason}")

    print(f"Query {query_execution_id} completed successfully.")
    return query_execution_id
```

```
In [3]: # set athena database name
athena_database = 'processed_data'
# set bucket used by Athena to store query results.
output_location = 's3://group9-ml-proj-athena-query-bucket-grw/' # change your initials
```

```
# create db query
create_db_query = f"CREATE DATABASE IF NOT EXISTS {athena_database};"
print("Creating database...")
run_athena_query(create_db_query, database='default', output_location=output_location)
```

Creating database...

Query submitted. Execution ID: 1f46bdd2-bf96-4e2d-a0b7-8802aa7ee779

Current query state: QUEUED

Current query state: SUCCEEDED

Query 1f46bdd2-bf96-4e2d-a0b7-8802aa7ee779 completed successfully.

Out[3]: '1f46bdd2-bf96-4e2d-a0b7-8802aa7ee779'

```
In [4]: # create external table
create_table_query = f"""
CREATE EXTERNAL TABLE IF NOT EXISTS {athena_database}.processed_data (
    FL_DATE STRING,
    CRS_DEP_TIME STRING,
    OP_UNIQUE_CARRIER STRING,
    DEST STRING,
    wind_dir_degrees DOUBLE,
    wind_speed_kt DOUBLE,
    wind_gust_kt DOUBLE,
    visibility_statute_mi DOUBLE,
    temperature_c DOUBLE,
    dewpoint_c DOUBLE,
    altimeter_hpa DOUBLE,
    Flight_Status STRING,
    Carrier_AA BOOLEAN,
    Carrier_AS BOOLEAN,
    Carrier_B6 BOOLEAN,
    Carrier_DL BOOLEAN,
    Carrier_F9 BOOLEAN,
    Carrier_G4 BOOLEAN,
    Carrier_HA BOOLEAN,
    Carrier_NK BOOLEAN,
    Carrier_OO BOOLEAN,
    Carrier_UA BOOLEAN,
    Carrier_WN BOOLEAN,
    DEST_REGION STRING,
    Dest_Region_Midwest BOOLEAN,
    Dest_Region_Mountain BOOLEAN,
    Dest_Region_Northeast BOOLEAN,
    Dest_Region_OCONUS BOOLEAN,
    Dest_Region_South BOOLEAN,
    Dest_Region_Southwest BOOLEAN,
```

```

        Dest_Region_West BOOLEAN,
        Flight_Status_Binary BOOLEAN
    )
    ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde'
    WITH SERDEPROPERTIES (
        'separatorChar' = ',',
        'quoteChar' = '\"'
    )
    LOCATION 's3://group9-ml-proj-processed-data-bucket-grw/all_data/'
    TBLPROPERTIES ('skip.header.line.count'='1');
    """
    print("Creating external table...")
    run_athena_query(create_table_query, database=athena_database, output_location=output_location)

    print("Athena database and table setup completed successfully.")

```

Creating external table...

Query submitted. Execution ID: e4362e9e-e2bb-4220-a981-aa1d309964aa

Current query state: QUEUED

Current query state: SUCCEEDED

Query e4362e9e-e2bb-4220-a981-aa1d309964aa completed successfully.

Athena database and table setup completed successfully.

```

In [1]: %%html
<p><b>Shutting down your kernel for this notebook to release resources.</b></p>
<button class="sm-command-button" data-commandlinker-command="kernelmenu:shutdown" style="display:none;">Shutdown Kernel</button>

<script>
try {
    els = document.getElementsByClassName("sm-command-button");
    els[0].click();
}
catch(err) {
    // NoOp
}
</script>

```

Shutting down your kernel for this notebook to release resources.

In []:

Weather Cleaning and Organization Notebook

This notebook will focus on converting the raw METAR and SPECI weather data into one usable data frame and exporting it to a CSV file that will be used to perform predictive analysis for on time departures out of San Diego International Airport. This finalized notebook will be used to perform a merge with the on time flight dataset resulting in weather data points for each of the flight information that was obtained.

```
In [1]: # Import libraries
import pandas as pd
import re

from datetime import datetime
from metar import Metar
```

```
In [2]: # Import the data files
df1 = pd.read_csv('raw_weather/KSAN_2023_wx.csv')
df2 = pd.read_csv('raw_weather/KSAN_2024_wx.csv')
df3 = pd.read_csv('raw_weather/KSAN_2025_wx.csv')

# create list of dfs
dfs = [df1, df2, df3]
```

```
/tmp/ipykernel_5660/2189777474.py:2: DtypeWarning: Columns (34,35,36,42,46,50,79) have mixed types. Specify dtype option on import or set low_memory=False.
  df1 = pd.read_csv('raw_weather/KSAN_2023_wx.csv')
/tmp/ipykernel_5660/2189777474.py:3: DtypeWarning: Columns (2,36,41) have mixed types. Specify dtype option on import or set low_memory=False.
  df2 = pd.read_csv('raw_weather/KSAN_2024_wx.csv')
```

```
In [3]: # Removing unnecessary rows besides REM which houses the raw METAR data and concatenating the dataframes
metar_dfs = []

for i, df in enumerate(dfs, start=1):
    if 'REM' in df.columns:
        metar_dfs.append(df[['REM']])
    else:
        print('REM column not found in df{i}')

# concatenating the dfs
comb_metar_df = pd.concat(metar_dfs, ignore_index=True)
```

```
In [4]: comb_metar_df.head()
```

Out[4]:

REM

```
0   SYN08672290 12366 8//06 10161 20128 30148 4016...
1   MET10512/31/22 16:10:03 SPECI KSAN 010010Z 180...
2   MET12712/31/22 16:51:03 METAR KSAN 010051Z 160...
3   MET12612/31/22 17:51:03 METAR KSAN 010151Z 150...
4   MET10512/31/22 18:16:03 SPECI KSAN 010216Z 140...
```

In [5]: `comb_metar_df.tail()`

Out[5]:

REM

```
30326   MET14103/06/25 21:01:01 SPECI KSAN 070501Z 330...
30327   MET13803/06/25 21:13:01 SPECI KSAN 070513Z 350...
30328   MET16903/06/25 21:51:01 METAR KSAN 070551Z 350...
30329   MET11903/06/25 22:51:01 METAR KSAN 070651Z 330...
30330   SOD77324 HR PRECIPITATION (IN): 0.59 WEATHER(C...
```

In [6]: *# Filter so that the data only contains METAR data and not extra synthetic information that does not conform to the standard METAR*
`comb_metar_df = comb_metar_df[comb_metar_df['REM'].str.startswith('MET', na=False)]`

In [7]: *# reset the index of the dataframe*
`comb_metar_df.reset_index(drop=True, inplace=True)`

In [8]: `print(comb_metar_df.head())`
`print(comb_metar_df.tail())`

```

                                REM
0  MET10512/31/22 16:10:03 SPECI KSAN 010010Z 180...
1  MET12712/31/22 16:51:03 METAR KSAN 010051Z 160...
2  MET12612/31/22 17:51:03 METAR KSAN 010151Z 150...
3  MET10512/31/22 18:16:03 SPECI KSAN 010216Z 140...
4  MET10812/31/22 18:38:03 SPECI KSAN 010238Z 140...

                                REM
24583 MET19203/06/25 20:51:01 METAR KSAN 070451Z 350...
24584 MET14103/06/25 21:01:01 SPECI KSAN 070501Z 330...
24585 MET13803/06/25 21:13:01 SPECI KSAN 070513Z 350...
24586 MET16903/06/25 21:51:01 METAR KSAN 070551Z 350...
24587 MET11903/06/25 22:51:01 METAR KSAN 070651Z 330...

```

In [9]: *# establish a prefix pattern to extract the information before the actual beginning of the METAR data*

```

prefix_pattern = re.compile(
    # the MET pattern and 3 digit code that we do not need
    r'^MET\d{3}'
    # local month
    r'(?P<month>\d{2})/'
    # local day
    r'(?P<day>\d{2})/'
    # local year
    r'(?P<year>\d{2})\s+'
    # local hour
    r'(?P<hour>\d{2}):'
    # local minute
    r'(?P<minute>\d{2}):'
    # local second
    r'(?P<second>\d{2})\s+'
    # rest of the actual metar
    r'(?P<metar>.*)$'
)

```

In [10]: *# create function that will parse the prefix data, dropping the MET and the three digit code but making a date and time column*

```

def parse_prefix(line):
    match = prefix_pattern.match(line)
    if not match:
        return None

    local_date = f"{match.group('month')}/{match.group('day')}/{match.group('year')}"
    local_time = f"{match.group('hour')}:{match.group('minute')}:{match.group('second')}"
    metar = match.group('metar')

    return {

```

```

        'local_date': local_date,
        'local_time': local_time,
        'metar': metar
    }

```

In [11]: *# create function that will parse the raw METAR data.*

```

def parse_metar(metar):
    try:
        parsed = Metar.Metar(metar)

        return {
            'station_id': parsed.station_id,
            'wind_dir_degrees': parsed.wind_dir.value() if parsed.wind_dir else None,
            'wind_speed_kt': parsed.wind_speed.value() if parsed.wind_speed else None,
            'wind_gust_kt': parsed.wind_gust.value() if parsed.wind_gust else None,
            'visibility_statute_mi': parsed.vis.value() if parsed.vis else None,
            'temperature_c': parsed.temp.value() if parsed.temp else None,
            'dewpoint_c': parsed.dewpt.value() if parsed.dewpt else None,
            'altimeter_hpa': parsed.press.value() if parsed.press else None,
        }
    except Exception as e:
        return None

```

In [12]: *# tie all the functions together in order to break up each field of every METAR observation into their own columns*

```

parsed_metar_data = []

# process each line of the comb_metar_df
for line in comb_metar_df['REM']:
    prefix_data = parse_prefix(line)
    if prefix_data is None:
        continue

    metar_data = parse_metar(prefix_data['metar'])
    if metar_data is None:
        continue

    combined_record = {**prefix_data, **metar_data}
    parsed_metar_data.append(combined_record)

# create the new df from parsed records
organized_metar_df = pd.DataFrame(parsed_metar_data)
organized_metar_df.reset_index(drop=True, inplace=True)

```



```
In [13]: print(organized_metar_df.head())
```

	local_date	local_time	metar	\
0	12/31/22	16:10:03	SPECI KSAN 010010Z 18004KT 10SM SCT009 BKN026 ...	
1	12/31/22	16:51:03	METAR KSAN 010051Z 16006KT 8SM -RA SCT008 BKN0...	
2	12/31/22	17:51:03	METAR KSAN 010151Z 15006KT 10SM FEW008 BKN021 ...	
3	12/31/22	18:16:03	SPECI KSAN 010216Z 14010KT 10SM FEW009 BKN031 ...	
4	12/31/22	18:38:03	SPECI KSAN 010238Z 14010G21KT 10SM FEW009 BKN0...	

	station_id	wind_dir_degrees	wind_speed_kt	wind_gust_kt	\
0	KSAN	180.0	4.0	NaN	
1	KSAN	160.0	6.0	NaN	
2	KSAN	150.0	6.0	NaN	
3	KSAN	140.0	10.0	NaN	
4	KSAN	140.0	10.0	21.0	

	visibility_statute_mi	temperature_c	dewpoint_c	altimeter_hpa
0	10.0	15.6	12.8	30.01
1	8.0	15.6	13.3	30.00
2	10.0	15.6	12.8	29.98
3	10.0	15.6	12.8	29.97
4	10.0	15.6	12.8	29.96

```
In [14]: # drop the metar column
organized_metar_df.drop(columns=['metar'], inplace=True)
print(organized_metar_df.head())
```

	local_date	local_time	station_id	wind_dir_degrees	wind_speed_kt	\
0	12/31/22	16:10:03	KSAN	180.0	4.0	
1	12/31/22	16:51:03	KSAN	160.0	6.0	
2	12/31/22	17:51:03	KSAN	150.0	6.0	
3	12/31/22	18:16:03	KSAN	140.0	10.0	
4	12/31/22	18:38:03	KSAN	140.0	10.0	

	wind_gust_kt	visibility_statute_mi	temperature_c	dewpoint_c	\
0	NaN	10.0	15.6	12.8	
1	NaN	8.0	15.6	13.3	
2	NaN	10.0	15.6	12.8	
3	NaN	10.0	15.6	12.8	
4	21.0	10.0	15.6	12.8	

	altimeter_hpa
0	30.01
1	30.00
2	29.98
3	29.97
4	29.96

```
In [15]: # replace all NaN values in the wind_gust_kt column with 0
organized_metar_df.fillna({'wind_gust_kt':0}, inplace=True)
print(organized_metar_df.head())
```

	local_date	local_time	station_id	wind_dir_degrees	wind_speed_kt	\
0	12/31/22	16:10:03	KSAN	180.0	4.0	
1	12/31/22	16:51:03	KSAN	160.0	6.0	
2	12/31/22	17:51:03	KSAN	150.0	6.0	
3	12/31/22	18:16:03	KSAN	140.0	10.0	
4	12/31/22	18:38:03	KSAN	140.0	10.0	

	wind_gust_kt	visibility_statute_mi	temperature_c	dewpoint_c	\
0	0.0	10.0	15.6	12.8	
1	0.0	8.0	15.6	13.3	
2	0.0	10.0	15.6	12.8	
3	0.0	10.0	15.6	12.8	
4	21.0	10.0	15.6	12.8	

	altimeter_hpa
0	30.01
1	30.00
2	29.98
3	29.97
4	29.96

```
In [16]: # output the data frame to a csv
organized_metar_df.to_csv('raw_weather/organized_metar_data.csv', index=False)
```

```
In [17]: %%html
<p><b>Shutting down your kernel for this notebook to release resources.</b></p>
<button class="sm-command-button" data-commandlinker-command="kernelmenu:shutdown" style="display:none;">Shutdown Kernel</button>

<script>
try {
  els = document.getElementsByClassName("sm-command-button");
  els[0].click();
}
catch(err) {
  // NoOp
}
</script>
```

Shutting down your kernel for this notebook to release resources.

```
In [ ]:
```

Raw Weather S3 Bucket Upload

The purpose of this notebook is to upload the rawest organized metar data into the raw_data S3 bucket. We will map this to an AWS Glue database in a later notebook.

```
In [1]: # import libraries
import boto3
import os
```

```
In [2]: # examine list of buckets available
s3 = boto3.client('s3')
response = s3.list_buckets()
print("Existing buckets:")
for bucket in response['Buckets']:
    print(bucket['Name'])
```

Existing buckets:

```
group9-ml-proj-model-artifacts-bucket-grw
group9-ml-proj-monitoring-bucket-grw
group9-ml-proj-processed-data-bucket-grw
group9-ml-proj-raw-data-bucket-grw
group9-ml-proj-training-artifacts-bucket-grw
sagemaker-studio-7sb8gxpljq9
sagemaker-studio-chxawrkbs1q
sagemaker-studio-o4i4ukj5y3b
sagemaker-us-east-1-804823187130
```

```
In [3]: # set local path of organized weather CSV file
local_file_path = 'raw_weather/organized_metar_data.csv'
```

```
In [4]: # set the raw data bucket variable
raw_bucket = "group9-ml-proj-raw-data-bucket-grw"
```

```
In [5]: # set the S3 key path, we are using flat structure because this will only hold three files
s3_key = "data/organized_metar_data.csv"
```

```
In [6]: # upload file to the S3 bucket
try:
    s3.upload_file(local_file_path, raw_bucket, s3_key)
    print(f"Uploaded {local_file_path} to s3://{raw_bucket}/{s3_key}")
```

```
except Exception as e:
    print("Error uploading file to S3:", e)
```

Uploaded raw_weather/organized_metar_data.csv to s3://group9-ml-proj-raw-data-bucket-grw/data/organized_metar_data.csv

Shutdown the kernel

```
In [1]: %%html
<p><b>Shutting down your kernel for this notebook to release resources.</b></p>
<button class="sm-command-button" data-commandlinker-command="kernelmenu:shutdown" style="display:none;">Shutdown Kernel</button>

<script>
try {
    els = document.getElementsByClassName("sm-command-button");
    els[0].click();
}
catch(err) {
    // NoOp
}
</script>
```

Shutting down your kernel for this notebook to release resources.

In []:

Flight Cleaning Notebook

This notebook will be responsible for merging all the flight information for the San Diego International Airport's On Time performance covering two years.

```
In [1]: # import libraries
import pandas as pd
```

```
In [2]: # import the csv files
Jan23 = pd.read_csv('raw_on-time/Jan23_OT_report.csv')
Feb23 = pd.read_csv('raw_on-time/Feb23_OT_report.csv')
Mar23 = pd.read_csv('raw_on-time/Mar23_OT_report.csv')
Apr23 = pd.read_csv('raw_on-time/Apr23_OT_report.csv')
May23 = pd.read_csv('raw_on-time/May23_OT_report.csv')
Jun23 = pd.read_csv('raw_on-time/Jun23_OT_report.csv')
Jul23 = pd.read_csv('raw_on-time/Jul23_OT_report.csv')
Aug23 = pd.read_csv('raw_on-time/Aug23_OT_report.csv')
Sep23 = pd.read_csv('raw_on-time/Sep23_OT_report.csv')
Oct23 = pd.read_csv('raw_on-time/Oct23_OT_report.csv')
Nov23 = pd.read_csv('raw_on-time/Nov23_OT_report.csv')
Dec23 = pd.read_csv('raw_on-time/Dec23_OT_report.csv')
Jan24 = pd.read_csv('raw_on-time/Jan24_OT_report.csv')
Feb24 = pd.read_csv('raw_on-time/Feb24_OT_report.csv')
Mar24 = pd.read_csv('raw_on-time/Mar24_OT_report.csv')
Apr24 = pd.read_csv('raw_on-time/Apr24_OT_report.csv')
May24 = pd.read_csv('raw_on-time/May24_OT_report.csv')
Jun24 = pd.read_csv('raw_on-time/Jun24_OT_report.csv')
Jul24 = pd.read_csv('raw_on-time/Jul24_OT_report.csv')
Aug24 = pd.read_csv('raw_on-time/Aug24_OT_report.csv')
Sep24 = pd.read_csv('raw_on-time/Sep24_OT_report.csv')
Oct24 = pd.read_csv('raw_on-time/Oct24_OT_report.csv')
Nov24 = pd.read_csv('raw_on-time/Nov24_OT_report.csv')
Dec24 = pd.read_csv('raw_on-time/Dec24_OT_report.csv')
```

```
In [3]: # List the df's in chronological order
dfs = [Jan23, Feb23, Mar23, Apr23, May23, Jun23, Jul23, Aug23, Sep23, Oct23, Nov23, Dec23,
        Jan24, Feb24, Mar24, Apr24, May24, Jun24, Jul24, Aug24, Sep24, Oct24, Nov24, Dec24]
```

```
In [4]: # merge df's
merged_df = pd.concat(dfs, ignore_index=True)
```

```
In [5]: # Filter all df so that only flights leaving San Diego are represented
SAN_df = merged_df.loc[merged_df['ORIGIN'] == 'SAN']
# Reset the index
SAN_df = SAN_df.reset_index(drop=True)
```

```
In [6]: # Get shape of the df
SAN_df.shape
```

```
Out[6]: (185600, 18)
```

```
In [7]: # get the data types of the df
SAN_df.dtypes
```

```
Out[7]: FL_DATE          object
OP_UNIQUE_CARRIER      object
OP_CARRIER_FL_NUM      int64
ORIGIN                  object
DEST                   object
CRS_DEP_TIME            int64
DEP_TIME               float64
DEP_DEL15              float64
DEP_DELAY_GROUP         float64
DEP_TIME_BLK           object
CANCELLED              float64
CANCELLATION_CODE       object
CARRIER_DELAY          float64
WEATHER_DELAY           float64
NAS_DELAY              float64
SECURITY_DELAY          float64
LATE_AIRCRAFT_DELAY     float64
DEP_DELAY              float64
dtype: object
```

```
In [8]: print(SAN_df.head())
```

	FL_DATE	OP_UNIQUE_CARRIER	OP_CARRIER_FL_NUM	ORIGIN	DEST	\
0	1/1/2023 12:00:00 AM	AA	1055	SAN	DFW	
1	1/1/2023 12:00:00 AM	AA	1651	SAN	CLT	
2	1/1/2023 12:00:00 AM	AA	1663	SAN	ORD	
3	1/1/2023 12:00:00 AM	AA	1765	SAN	DFW	
4	1/1/2023 12:00:00 AM	AA	1939	SAN	DFW	

	CRS_DEP_TIME	DEP_TIME	DEP_DEL15	DEP_DELAY_GROUP	DEP_TIME_BLK	CANCELLED	\
0	715	713.0	0.0	-1.0	0700-0759	0.0	
1	615	644.0	1.0	1.0	0600-0659	0.0	
2	2259	2250.0	0.0	-1.0	2200-2259	0.0	
3	815	810.0	0.0	-1.0	0800-0859	0.0	
4	1450	1447.0	0.0	-1.0	1400-1459	0.0	

	CANCELLATION_CODE	CARRIER_DELAY	WEATHER_DELAY	NAS_DELAY	SECURITY_DELAY	\
0	NaN	NaN	NaN	NaN	NaN	
1	NaN	NaN	NaN	NaN	NaN	
2	NaN	NaN	NaN	NaN	NaN	
3	NaN	NaN	NaN	NaN	NaN	
4	NaN	NaN	NaN	NaN	NaN	

	LATE_AIRCRAFT_DELAY	DEP_DELAY
0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	NaN
4	NaN	NaN

```
In [9]: # adjust the FL_DATE column to datetime
SAN_df['FL_DATE'] = pd.to_datetime(SAN_df['FL_DATE']).dt.date
```

/tmp/ipykernel_5818/1251424886.py:2: UserWarning: Could not infer format, so each element will be parsed individually, falling back to `dateutil`. To ensure parsing is consistent and as-expected, please specify a format.

```
SAN_df['FL_DATE'] = pd.to_datetime(SAN_df['FL_DATE']).dt.date
```

```
In [10]: # convert CRS_DEP_TIME to a string
SAN_df['CRS_DEP_TIME'] = SAN_df['CRS_DEP_TIME'].astype(str).str.zfill(4)
# convert CRS_DEP_TIME to time format
SAN_df['CRS_DEP_TIME'] = pd.to_datetime(SAN_df['CRS_DEP_TIME'], format='%H%M').dt.time
```

```
In [11]: print(SAN_df.head())
```


	FL_DATE	OP_UNIQUE_CARRIER	OP_CARRIER_FL_NUM	ORIGIN	DEST	CRS_DEP_TIME \
0	2023-01-01	AA	1055	SAN	DFW	07:15:00
1	2023-01-01	AA	1651	SAN	CLT	06:15:00
2	2023-01-01	AA	1663	SAN	ORD	22:59:00
3	2023-01-01	AA	1765	SAN	DFW	08:15:00
4	2023-01-01	AA	1939	SAN	DFW	14:50:00

	DEP_TIME	DEP_DEL15	DEP_DELAY_GROUP	DEP_TIME_BLK	CANCELLED \
0	713.0	0.0	-1.0	0700-0759	0.0
1	644.0	1.0	1.0	0600-0659	0.0
2	2250.0	0.0	-1.0	2200-2259	0.0
3	810.0	0.0	-1.0	0800-0859	0.0
4	1447.0	0.0	-1.0	1400-1459	0.0

	CANCELLATION_CODE	CARRIER_DELAY	WEATHER_DELAY	NAS_DELAY	SECURITY_DELAY \
0	NaN	NaN	NaN	NaN	NaN
1	NaN	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN	NaN

	LATE_AIRCRAFT_DELAY	DEP_DELAY
0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	NaN
4	NaN	NaN

In [12]: *# create function to handle missing DEP_TIME for cancelled flights*

```
def convert_dep_time(val):
    if pd.isnull(val):
        return None
    try:
        val_float = float(val)
        val_int = int(val_float)
        time_str = str(val_int).zfill(4)
        return pd.to_datetime(time_str, format='%H%M').time()
    except Exception:
        return None
```

In [13]: *# apply the function to the DEP_TIME column*

```
SAN_df['DEP_TIME'] = SAN_df['DEP_TIME'].apply(convert_dep_time)
```

In [14]: `print(SAN_df.head())`

	FL_DATE	OP_UNIQUE_CARRIER	OP_CARRIER_FL_NUM	ORIGIN	DEST	CRS_DEP_TIME	\
0	2023-01-01	AA	1055	SAN	DFW	07:15:00	
1	2023-01-01	AA	1651	SAN	CLT	06:15:00	
2	2023-01-01	AA	1663	SAN	ORD	22:59:00	
3	2023-01-01	AA	1765	SAN	DFW	08:15:00	
4	2023-01-01	AA	1939	SAN	DFW	14:50:00	

	DEP_TIME	DEP_DEL15	DEP_DELAY_GROUP	DEP_TIME_BLK	CANCELLED	\
0	07:13:00	0.0	-1.0	0700-0759	0.0	
1	06:44:00	1.0	1.0	0600-0659	0.0	
2	22:50:00	0.0	-1.0	2200-2259	0.0	
3	08:10:00	0.0	-1.0	0800-0859	0.0	
4	14:47:00	0.0	-1.0	1400-1459	0.0	

	CANCELLATION_CODE	CARRIER_DELAY	WEATHER_DELAY	NAS_DELAY	SECURITY_DELAY	\
0	NaN	NaN	NaN	NaN	NaN	
1	NaN	NaN	NaN	NaN	NaN	
2	NaN	NaN	NaN	NaN	NaN	
3	NaN	NaN	NaN	NaN	NaN	
4	NaN	NaN	NaN	NaN	NaN	

	LATE_AIRCRAFT_DELAY	DEP_DELAY
0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	NaN
4	NaN	NaN

```
In [15]: # convert the DEP_DEL15 column to int
SAN_df['DEP_DEL15'] = SAN_df['DEP_DEL15'].fillna(0).astype(int)
```

```
In [16]: # convert CANCELLED column to int
SAN_df['CANCELLED'] = SAN_df['CANCELLED'].fillna(0).astype(int)
```

```
In [17]: # check what values are present in cancellation_code
SAN_df['CANCELLATION_CODE'].value_counts()
```

```
Out[17]: CANCELLATION_CODE
A      992
B      902
C      222
D         2
Name: count, dtype: int64
```

```
In [18]: # fill missing values with None in CANCELLATION_CODE
SAN_df['CANCELLATION_CODE'] = SAN_df['CANCELLATION_CODE'].fillna('None')
```

```
In [19]: # check what values are present in CARRIER_DELAY
SAN_df['CARRIER_DELAY'].value_counts()
```

```
Out[19]: CARRIER_DELAY
0.0      23180
1.0       887
2.0       871
3.0       776
4.0       722
...
263.0      1
1055.0      1
436.0      1
863.0      1
284.0      1
Name: count, Length: 546, dtype: int64
```

```
In [20]: # convert the CARRIER_DELAY column to int
SAN_df['CARRIER_DELAY'] = SAN_df['CARRIER_DELAY'].fillna(0).astype(int)
```

```
In [21]: # convert the WEATHER_DELAY column to int
SAN_df['WEATHER_DELAY'] = SAN_df['WEATHER_DELAY'].fillna(0).astype(int)
```

```
In [22]: # convert the NAS_DELAY column to int
SAN_df['NAS_DELAY'] = SAN_df['NAS_DELAY'].fillna(0).astype(int)
```

```
In [23]: # convert the SECURITY_DELAY column to int
SAN_df['SECURITY_DELAY'] = SAN_df['SECURITY_DELAY'].fillna(0).astype(int)
```

```
In [24]: # convert the LATE_AIRCRAFT_DELAY column to int
SAN_df['LATE_AIRCRAFT_DELAY'] = SAN_df['LATE_AIRCRAFT_DELAY'].fillna(0).astype(int)
```

```
In [25]: # reorder the columns so times are next to the date
SAN_df = SAN_df[['FL_DATE', 'CRS_DEP_TIME', 'DEP_TIME', 'OP_UNIQUE_CARRIER', 'OP_CARRIER_FL_NUM', 'ORIGIN', 'DEST', 'DEP_DEL15',
print(SAN_df.head())
```

	FL_DATE	CRS_DEP_TIME	DEP_TIME	OP_UNIQUE_CARRIER	OP_CARRIER_FL_NUM	\
0	2023-01-01	07:15:00	07:13:00	AA	1055	
1	2023-01-01	06:15:00	06:44:00	AA	1651	
2	2023-01-01	22:59:00	22:50:00	AA	1663	
3	2023-01-01	08:15:00	08:10:00	AA	1765	
4	2023-01-01	14:50:00	14:47:00	AA	1939	

	ORIGIN	DEST	DEP_DEL15	CANCELLED	CANCELLATION_CODE	CARRIER_DELAY	\
0	SAN	DFW	0	0	None	0	
1	SAN	CLT	1	0	None	0	
2	SAN	ORD	0	0	None	0	
3	SAN	DFW	0	0	None	0	
4	SAN	DFW	0	0	None	0	

	WEATHER_DELAY	NAS_DELAY	SECURITY_DELAY	LATE_AIRCRAFT_DELAY
0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0

```
In [26]: # Save the df to a csv file
SAN_df.to_csv('raw_on-time/SAN_OT_report.csv', index=False)
```

```
In [27]: %%html
<p><b>Shutting down your kernel for this notebook to release resources.</b></p>
<button class="sm-command-button" data-commandlinker-command="kernelmenu:shutdown" style="display:none;">Shutdown Kernel</button>

<script>
try {
    els = document.getElementsByClassName("sm-command-button");
    els[0].click();
}
catch(err) {
    // NoOp
}
</script>
```

Shutting down your kernel for this notebook to release resources.

Raw Flight Data Merging and Uploading to S3

```
In [1]: # import libraries
import boto3
import os
```

```
In [2]: # examine list of buckets available
s3 = boto3.client('s3')
response = s3.list_buckets()
print("Existing buckets:")
for bucket in response['Buckets']:
    print(bucket['Name'])
```

Existing buckets:

```
group9-ml-proj-model-artifacts-bucket-grw
group9-ml-proj-monitoring-bucket-grw
group9-ml-proj-processed-data-bucket-grw
group9-ml-proj-raw-data-bucket-grw
group9-ml-proj-training-artifacts-bucket-grw
sagemaker-studio-7sb8gxpljq9
sagemaker-studio-chxawrkbs1q
sagemaker-studio-o4i4ukj5y3b
sagemaker-us-east-1-804823187130
```

```
In [3]: # set local path of merged CSV file
local_file_path = 'raw_on-time/SAN_OT_report.csv'
```

```
In [4]: # set the raw data bucket variable
raw_bucket = "group9-ml-proj-raw-data-bucket-grw"
```

```
In [5]: # set the S3 key path, we are using flat structure because this will only hold two files
s3_key = "data/SAN_OT_report.csv"
```

```
In [6]: # upload file to the S3 bucket
try:
    s3.upload_file(local_file_path, raw_bucket, s3_key)
    print(f"Uploaded {local_file_path} to s3://{raw_bucket}/{s3_key}")
except Exception as e:
    print("Error uploading file to S3:", e)
```

Uploaded raw_on-time/SAN_OT_report.csv to s3://group9-ml-proj-raw-data-bucket-grw/data/SAN_OT_report.csv

Shutdown the kernel

```
In [1]: %%html
<p><b>Shutting down your kernel for this notebook to release resources.</b></p>
<button class="sm-command-button" data-commandlinker-command="kernelmenu:shutdown" style="display:none;">Shutdown Kernel</button>

<script>
try {
  els = document.getElementsByClassName("sm-command-button");
  els[0].click();
}
catch(err) {
  // NoOp
}
</script>
```

Shutting down your kernel for this notebook to release resources.

In []:

Merged Notebook

This notebook will serve the purpose of merging the organized METAR and San Diego On Time report into one usable dataset.

The merge will focus on merging on the variables local_date and local_time from the organized_metar_data.csv file and flight_date and CRS_DEP_TIME from the SAN_OT_report.csv file. In the event that the local_time and CRS_DEP_TIME do not exactly match the closest weather observation prior to CRS_DEP_TIME will be used.

```
In [2]: # import libraries
import pandas as pd
```

```
In [3]: # read in the data
wx_data = pd.read_csv("/home/sagemaker-user/ADS-508-Project/2_raw_data/raw_weather/organized_metar_data.csv")
flt_data = pd.read_csv("/home/sagemaker-user/ADS-508-Project/2_raw_data/raw_on-time/SAN_OT_report.csv")
```

```
In [4]: # create a datetime column for the weather data
wx_data['wx_datetime'] = pd.to_datetime(wx_data['local_date'] + ' ' + wx_data['local_time'],
                                       format="%m/%d/%y %H:%M:%S",
                                       errors='coerce')

wx_data[['local_date', 'local_time', 'wx_datetime']].head()
```

```
Out[4]:
```

	local_date	local_time	wx_datetime
0	12/31/22	16:10:03	2022-12-31 16:10:03
1	12/31/22	16:51:03	2022-12-31 16:51:03
2	12/31/22	17:51:03	2022-12-31 17:51:03
3	12/31/22	18:16:03	2022-12-31 18:16:03
4	12/31/22	18:38:03	2022-12-31 18:38:03

```
In [5]: # create a datetime column for the flight data
flt_data['flt_datetime'] = pd.to_datetime(flt_data['FL_DATE'] + ' ' + flt_data['CRS_DEP_TIME'],
                                       format="%Y-%m-%d %H:%M:%S",
                                       errors='coerce')

flt_data[['FL_DATE', 'CRS_DEP_TIME', 'flt_datetime']].head()
```

Out[5]:

	FL_DATE	CRS_DEP_TIME	flt_datetime
0	2023-01-01	07:15:00	2023-01-01 07:15:00
1	2023-01-01	06:15:00	2023-01-01 06:15:00
2	2023-01-01	22:59:00	2023-01-01 22:59:00
3	2023-01-01	08:15:00	2023-01-01 08:15:00
4	2023-01-01	14:50:00	2023-01-01 14:50:00

```
In [6]: # sort the flight data by flt_datetime
flt_data = flt_data.sort_values(by='flt_datetime')
```

```
In [7]: # check both sorted by created datetime fields
print("METAR is monotonic:", wx_data["wx_datetime"].is_monotonic_increasing)
print("Flights are monotonic:", flt_data["flt_datetime"].is_monotonic_increasing)
```

```
METAR is monotonic: True
Flights are monotonic: True
```

```
In [8]: # merge the dataframes on the datetime fields created prior
merged_data = pd.merge_asof(flt_data,
                             wx_data, left_on='flt_datetime',
                             right_on='wx_datetime',
                             direction='backward')

merged_data.head()
```



```
Out[8]:
```

	FL_DATE	CRS_DEP_TIME	DEP_TIME	OP_UNIQUE_CARRIER	OP_CARRIER_FL_NUM	ORIGIN	DEST	DEP_DEL15	CANCELLED	CANCELLATION
0	2023-01-01	06:10:00	06:05:00	DL	977	SAN	DTW	0	0	
1	2023-01-01	06:15:00	06:17:00	DL	2423	SAN	SLC	0	0	
2	2023-01-01	06:15:00	06:44:00	AA	1651	SAN	CLT	1	0	
3	2023-01-01	06:15:00	06:15:00	DL	820	SAN	ATL	0	0	
4	2023-01-01	06:15:00	06:13:00	DL	914	SAN	MSP	0	0	

5 rows × 27 columns



```
In [9]: # check the column names
merged_data.columns
```

```
Out[9]: Index(['FL_DATE', 'CRS_DEP_TIME', 'DEP_TIME', 'OP_UNIQUE_CARRIER',
              'OP_CARRIER_FL_NUM', 'ORIGIN', 'DEST', 'DEP_DEL15', 'CANCELLED',
              'CANCELLATION_CODE', 'CARRIER_DELAY', 'WEATHER_DELAY', 'NAS_DELAY',
              'SECURITY_DELAY', 'LATE_AIRCRAFT_DELAY', 'flt_datetime', 'local_date',
              'local_time', 'station_id', 'wind_dir_degrees', 'wind_speed_kt',
              'wind_gust_kt', 'visibility_statute_mi', 'temperature_c', 'dewpoint_c',
              'altimeter_hpa', 'wx_datetime'],
              dtype='object')
```

```
In [10]: # drop unnecessary columns
merged_data = merged_data.drop(columns=['local_date', 'local_time', 'wx_datetime', 'station_id', 'flt_datetime', 'ORIGIN'])
```

```
In [11]: # save the merged data
merged_data.to_csv("/home/sagemaker-user/ADS-508-Project/2_raw_data/raw_merged/raw_merged_data.csv", index=False)
```

```
In [12]: %%html
<p><b>Shutting down your kernel for this notebook to release resources.</b></p>
<button class="sm-command-button" data-commandlinker-command="kernelmenu:shutdown" style="display:none;">Shutdown Kernel</button>

<script>
try {
    els = document.getElementsByClassName("sm-command-button");
```

```
    els[0].click();  
}  
catch(err) {  
    // NoOp  
}  
</script>
```

Shutting down your kernel for this notebook to release resources.

In []:

Upload Raw Merged Dataset to S3

```
In [1]: # import libraries
import boto3
import os
```

```
In [2]: # examine list of buckets available
s3 = boto3.client('s3')
response = s3.list_buckets()
print("Existing buckets:")
for bucket in response['Buckets']:
    print(bucket['Name'])
```

Existing buckets:

```
group9-ml-proj-model-artifacts-bucket-grw
group9-ml-proj-monitoring-bucket-grw
group9-ml-proj-processed-data-bucket-grw
group9-ml-proj-raw-data-bucket-grw
group9-ml-proj-training-artifacts-bucket-grw
sagemaker-studio-7sb8gxpljq9
sagemaker-studio-chxawrkbs1q
sagemaker-studio-o4i4ukj5y3b
sagemaker-us-east-1-804823187130
```

```
In [3]: # set local path of merged CSV file
local_file_path = 'raw_merged/raw_merged_data.csv'
```

```
In [4]: # set the raw data bucket variable
raw_bucket = "group9-ml-proj-raw-data-bucket-grw"
```

```
In [5]: # set the S3 key path, we are using flat structure because this will only hold two files
s3_key = "data/raw_merged_data.csv"
```

```
In [6]: # upload file to the S3 bucket
try:
    s3.upload_file(local_file_path, raw_bucket, s3_key)
    print(f"Uploaded {local_file_path} to s3://{raw_bucket}/{s3_key}")
except Exception as e:
    print("Error uploading file to S3:", e)
```

Uploaded raw_merged/raw_merged_data.csv to s3://group9-ml-proj-raw-data-bucket-grw/data/raw_merged_data.csv

```
In [1]: %%html
<p><b>Shutting down your kernel for this notebook to release resources.</b></p>
<button class="sm-command-button" data-commandlinker-command="kernelmenu:shutdown" style="display:none;">Shutdown Kernel</button>

<script>
try {
  els = document.getElementsByClassName("sm-command-button");
  els[0].click();
}
catch(err) {
  // NoOp
}
</script>
```

Shutting down your kernel for this notebook to release resources.

In []:

Exploratory Data Analysis Section

```
In [2]: # import libraries
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
import os
import boto3

from windrose import WindroseAxes
```

```
In [3]: # read in the csv data file into a pandas dataframe
raw_df = pd.read_csv("../2_raw_data/raw_merged/raw_merged_data.csv")
```

```
In [4]: # Get shape of the data
raw_df.shape
```

```
Out[4]: (185600, 21)
```

```
In [5]: # Get a list of the column names
raw_df.columns
```

```
Out[5]: Index(['FL_DATE', 'CRS_DEP_TIME', 'DEP_TIME', 'OP_UNIQUE_CARRIER',
              'OP_CARRIER_FL_NUM', 'DEST', 'DEP_DEL15', 'CANCELLED',
              'CANCELLATION_CODE', 'CARRIER_DELAY', 'WEATHER_DELAY', 'NAS_DELAY',
              'SECURITY_DELAY', 'LATE_AIRCRAFT_DELAY', 'wind_dir_degrees',
              'wind_speed_kt', 'wind_gust_kt', 'visibility_statute_mi',
              'temperature_c', 'dewpoint_c', 'altimeter_hpa'],
              dtype='object')
```

```
In [6]: # get information on the df
raw_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 185600 entries, 0 to 185599
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   FL_DATE                185600 non-null object
1   CRS_DEP_TIME           185600 non-null object
2   DEP_TIME               183548 non-null object
3   OP_UNIQUE_CARRIER     185600 non-null object
4   OP_CARRIER_FL_NUM     185600 non-null int64
5   DEST                  185600 non-null object
6   DEP_DEL15              185600 non-null int64
7   CANCELLED              185600 non-null int64
8   CANCELLATION_CODE      2118 non-null  object
9   CARRIER_DELAY         185600 non-null int64
10  WEATHER_DELAY          185600 non-null int64
11  NAS_DELAY              185600 non-null int64
12  SECURITY_DELAY         185600 non-null int64
13  LATE_AIRCRAFT_DELAY    185600 non-null int64
14  wind_dir_degrees       170031 non-null float64
15  wind_speed_kt          185600 non-null float64
16  wind_gust_kt           185600 non-null float64
17  visibility_statute_mi  185600 non-null float64
18  temperature_c          185600 non-null float64
19  dewpoint_c             185600 non-null float64
20  altimeter_hpa          185600 non-null float64
dtypes: float64(7), int64(8), object(6)
memory usage: 29.7+ MB

```

```

In [7]: # get number of nulls in each column
raw_df.isnull().sum()

```

```

Out[7]: FL_DATE          0
        CRS_DEP_TIME    0
        DEP_TIME        2052
        OP_UNIQUE_CARRIER 0
        OP_CARRIER_FL_NUM 0
        DEST            0
        DEP_DEL15        0
        CANCELLED        0
        CANCELLATION_CODE 183482
        CARRIER_DELAY    0
        WEATHER_DELAY     0
        NAS_DELAY         0
        SECURITY_DELAY    0
        LATE_AIRCRAFT_DELAY 0
        wind_dir_degrees 15569
        wind_speed_kt      0
        wind_gust_kt       0
        visibility_statute_mi 0
        temperature_c      0
        dewpoint_c         0
        altimeter_hpa      0
        dtype: int64

```

```

In [8]: # get summary statistics of the dataframe
raw_df.describe()

```

```

Out[8]:

```

	OP_CARRIER_FL_NUM	DEP_DEL15	CANCELLED	CARRIER_DELAY	WEATHER_DELAY	NAS_DELAY	SECURITY_DELAY	LATE_AIRCRAFT_DELAY
count	185600.000000	185600.000000	185600.000000	185600.000000	185600.000000	185600.000000	185600.000000	185600.000000
mean	1992.292236	0.095264	0.011412	3.595609	0.565388	3.030501	0.022328	0.022328
std	1281.470818	0.293580	0.106214	33.489037	12.527305	15.422100	1.125169	1.125169
min	6.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	894.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
50%	1939.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
75%	2924.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
max	8784.000000	1.000000	1.000000	3300.000000	1171.000000	1431.000000	276.000000	276.000000

Target Variable Exploration of Delayed and Cancelled flights

```
In [9]: # count the number of DEP_DEL15 values
raw_df['DEP_DEL15'].value_counts()
```

```
Out[9]: DEP_DEL15
0      167919
1       17681
Name: count, dtype: int64
```


```
In [10]: # get unique values of DEP_DEL15
raw_df['DEP_DEL15'].unique()
```

```
Out[10]: array([0, 1])
```


```
In [11]: # count the number of CANCELLED values
raw_df['CANCELLED'].value_counts()
```

```
Out[11]: CANCELLED
0      183482
1        2118
Name: count, dtype: int64
```

```
In [12]: # visualize the distribution of delayed flights
plt.figure(figsize=(10,6))
sns.countplot(data=raw_df, x='DEP_DEL15')
plt.title('Distribution of Delayed Flights')
plt.xlabel('Delayed (1 = Yes, 0 = No)')
plt.ylabel('Number of Flights')
plt.tight_layout()
plt.show()
```


```
In [13]: # visualize distribution of canceled flights
plt.figure(figsize=(10,6))
sns.countplot(data=raw_df, x='CANCELLED')
plt.title('Distribution of Canceled Flights')
plt.xlabel('Canceled (1 = Yes, 0 = No)')
plt.ylabel('Number of Flights')
plt.tight_layout()
plt.show()
```



```
In [14]: # explore the distribution of carriers and flights
# get count of flights by airline
carrier_counts = raw_df['OP_UNIQUE_CARRIER'].value_counts().reset_index()
carrier_counts.columns = ['Carrier', 'Count of Flights']

# plot the count of flights by airline
plt.figure(figsize=(10,6))
sns.barplot(data=carrier_counts, x='Carrier', y='Count of Flights')
plt.title('Count of Flights by Carrier')
plt.xlabel('Carrier')
plt.ylabel('Count of Flights')
```

```
plt.tight_layout()
plt.show()
```



```
In [15]: # create new category in dataframe for eda
raw_eda_df = raw_df.copy()
```

```
In [16]: """
Function that will categorize Canceled, Delayed, and On-Time flights
"""
def classify_fl_status(row):
    if row['CANCELLED'] == 1:
        return 'Cancelled'
```

```

elif row['DEP_DEL15'] == 1:
    return 'Delayed'
elif row['DEP_DEL15'] == 0:
    return 'On-Time'
else:
    return 'Unknown'

```

```

In [17]: # create the new column in the dataframe by applying the function
raw_eda_df['Flight_Status'] = raw_eda_df.apply(classify_fl_status, axis=1)


# get the counts of the flight status column
fl_status_counts = raw_eda_df['Flight_Status'].value_counts().reset_index()
# print(fl_status_counts)

# print the percentage of each flight status
fl_status_counts['Percentage'] = (fl_status_counts['count'] / fl_status_counts['count'].sum()) * 100
print(fl_status_counts)

# visualize the distribution of flight status
plt.figure(figsize=(10,6))
sns.countplot(data=raw_eda_df, x='Flight_Status')
plt.title('Distribution of Flight Status')
plt.xlabel('Flight Status')
plt.ylabel('Number of Flights')
plt.tight_layout()
plt.show()

```

	Flight_Status	count	Percentage
0	On-Time	165811	89.337823
1	Delayed	17671	9.521013
2	Cancelled	2118	1.141164



```
In [18]: # explore flight carriers and statuses
# get the count of flights by carrier and status
carrier_status_counts = raw_eda_df.groupby(['OP_UNIQUE_CARRIER', 'Flight_Status']).size().reset_index()
carrier_status_counts.columns = ['Carrier', 'Flight_Status', 'Count']


# pivot table for plotting
carrier_status_pivot = carrier_status_counts.pivot(index='Carrier', columns='Flight_Status', values='Count').fillna(0)
print(carrier_status_pivot)
# sort by total flight count per carrier
carrier_status_pivot['Total'] = carrier_status_pivot.sum(axis=1)
carrier_status_pivot = carrier_status_pivot.sort_values(by='Total', ascending=False)
carrier_status_pivot = carrier_status_pivot.drop(columns='Total')
```

```

# plot stacked bar chart
carrier_status_pivot.plot(kind='bar', stacked=True, figsize=(12, 6))
plt.title('Count of Flights by Carrier and Flight Status')
plt.xlabel('Carrier')
plt.ylabel('Number of Flights')
plt.legend(title='Flight Status')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

```

Carrier	Cancelled	Delayed	On-Time
AA	192	1483	16001
AS	326	1495	17581
B6	29	634	3060
DL	141	1211	17482
F9	133	826	4971
G4	9	158	425
HA	8	173	1253
NK	87	481	4263
OO	139	1453	20938
UA	320	1378	18735
WN	734	8379	61102



```
In [19]: # noticing mismatch in counts before and after
# count how many rows have both cancelled and delayed
both_cancelled_and_delayed = raw_eda_df[
    (raw_eda_df['DEP_DEL15'] == 1) &
    (raw_eda_df['CANCELLED'] == 1)
]

print("Rows where a cancelled flight was also marked as delayed:", len(both_cancelled_and_delayed))

## There are 10 flights that were classified as both cancelled and delayed,
## this accounts for why the counts seem to be off when we examined the counts of the columns previously
## these rows will be classified as cancelled in the processed dataset
```

Rows where a cancelled flight was also marked as delayed: 10

```
In [20]: # one hot encode the carrier column
carrier_dummies = pd.get_dummies(raw_eda_df['OP_UNIQUE_CARRIER'], prefix='Carrier')
```

```
# concatenate the carrier dummies with the original dataframe, we are not going to get rid of the extra column just yet.
raw_eda_df = pd.concat([raw_eda_df, carrier_dummies], axis=1)

print(raw_eda_df.head())
```

	FL_DATE	CRS_DEP_TIME	DEP_TIME	OP_UNIQUE_CARRIER	OP_CARRIER_FL_NUM	\
0	2023-01-01	06:10:00	06:05:00	DL	977	
1	2023-01-01	06:15:00	06:17:00	DL	2423	
2	2023-01-01	06:15:00	06:44:00	AA	1651	
3	2023-01-01	06:15:00	06:15:00	DL	820	
4	2023-01-01	06:15:00	06:13:00	DL	914	

	DEST	DEP_DEL15	CANCELLED	CANCELLATION_CODE	CARRIER_DELAY	...	\
0	DTW	0	0	NaN	0	...	
1	SLC	0	0	NaN	0	...	
2	CLT	1	0	NaN	0	...	
3	ATL	0	0	NaN	0	...	
4	MSP	0	0	NaN	0	...	

	Carrier_AS	Carrier_B6	Carrier_DL	Carrier_F9	Carrier_G4	Carrier_HA	\
0	False	False	True	False	False	False	
1	False	False	True	False	False	False	
2	False	False	False	False	False	False	
3	False	False	True	False	False	False	
4	False	False	True	False	False	False	

	Carrier_NK	Carrier_OO	Carrier_UA	Carrier_WN
0	False	False	False	False
1	False	False	False	False
2	False	False	False	False
3	False	False	False	False
4	False	False	False	False

[5 rows x 33 columns]

Exploration of Weather Features

```
In [21]: # get dtype of weather columns
weather_columns = ['wind_dir_degrees', 'wind_speed_kt', 'visibility_statute_mi', 'temperature_c', 'dewpoint_c', 'altimeter_hpa']
raw_eda_df[weather_columns].dtypes
```



```
Out[21]: wind_dir_degrees      float64
wind_speed_kt                float64
visibility_statute_mi        float64
temperature_c                float64
dewpoint_c                  float64
altimeter_hpa                float64
dtype: object
```

```
In [22]: # check unique values for wind direction
raw_eda_df['wind_dir_degrees'].unique()
```

```
Out[22]: array([290., 280., 270.,  0., 130., 140., nan, 210., 240., 260., 220.,
        160., 120., 150., 170., 190., 180., 200., 250., 350., 340., 320.,
        300., 310., 330., 20., 90., 50., 360., 10., 40., 230., 60.,
        70., 110., 100., 30., 80.])
```

```
In [23]: # since 0 and 360 are the same, we will replace 0 with 360
raw_eda_df['wind_dir_degrees'] = raw_eda_df['wind_dir_degrees'].replace(0.0, 360.0)

# Address missing values in wind_dir_degrees to change to integer and address missing values
# Missing values are attributed to what is known as Variable Wind speeds, which are often associated with relatively calm wind speeds
# We will replace these missing values with the mode of the wind_dir_degrees column
# get the mode of wind_dir_degrees
mode_wind_dir = raw_eda_df['wind_dir_degrees'].mode()[0]
print("Mode of wind_dir_degrees:", mode_wind_dir)

# replace missing values with the mode
raw_eda_df['wind_dir_degrees'] = raw_eda_df['wind_dir_degrees'].fillna(mode_wind_dir)

# get the count of missing values in wind_dir_degrees
missing_wind_dir = raw_eda_df['wind_dir_degrees'].isnull().sum()
print("Missing values in wind_dir_degrees:", missing_wind_dir)
```

Mode of wind_dir_degrees: 360.0

Missing values in wind_dir_degrees: 0

```
In [24]: # convert the wind_dir_degrees column to integer
raw_eda_df['wind_dir_degrees'] = raw_eda_df['wind_dir_degrees'].astype(int)
```

```
In [25]: # explore unique values for wind speed
print(raw_eda_df['wind_speed_kt'].unique())

# convert them to integers
raw_eda_df['wind_speed_kt'] = raw_eda_df['wind_speed_kt'].astype(int)
```

```
[ 7. 11. 17. 15. 13. 16. 20. 14. 19. 18. 12.  9.  8.  0.  3.  5.  6. 10.
 4. 21. 22.]
```

```
In [26]: # explore unique values for wind gust
print(raw_eda_df['wind_gust_kt'].unique())
```

```
# convert them to integers
raw_eda_df['wind_gust_kt'] = raw_eda_df['wind_gust_kt'].astype(int)
```

```
[19. 25.  0. 24. 23. 28. 26. 21. 18. 20. 22. 16. 17. 27. 15. 30. 34. 14.
 29. 42. 32. 35. 31. 33.]
```

```
In [27]: # explore unique values for visibility
print(raw_eda_df['visibility_statute_mi'].unique())
```

```
[ 6.    3.    9.    10.    8.    4.    2.5    1.5    7.    2.
 5.    1.    1.75  1.25  0.75  0.25  0.5    0.125]
```

```
In [28]: # check temperature unique values
print(raw_eda_df['temperature_c'].unique())
```

```
[ 13.3  13.9  14.4  15.   15.6  11.7  12.2  12.8  14.   16.1  16.   16.7
 17.8  17.2   9.4  10.   11.1   8.3   7.8  10.6  18.9  19.4  18.3   8.9
 20.6  21.1  22.2  20.   17.   13.    5.6   6.1   4.4   7.2   5.   21.7
  6.7  24.4  23.3  22.8  12.   23.9  25.6  26.7  26.1  27.2  25.   18.
  0.   19.   21.   27.8  27.   22.   23.   28.9  28.3 -17.8  24.   11.
  3.9  29.4  30.   30.6  31.1  31.7  32.2  32.8  34.4]
```

```
In [29]: # check dewpoint unique values
print(raw_eda_df['dewpoint_c'].unique())
```

```
[ 10.    9.4   8.3   8.9   7.8  10.6  11.   11.1  11.7  12.   12.2  13.3
 12.8  13.9  14.   13.   14.4   5.    5.6   6.1   3.3   4.4   2.8   0.6
  2.2   3.9   6.7   7.2   1.1   8.   -2.2  -1.7   1.7  -3.9  -3.3  -6.7
 -6.1  -4.4  -0.6   0.   -5.   -7.8  -7.2  -8.3  -5.6  -2.8  -1.1  -8.9
  7.  -10.6 -12.8 -12.2 -11.7 -13.3   9.    6.   15.   15.6  16.   16.1
 16.7  17.2  17.8  18.3  18.9  17.   18.   19.   19.4  21.1  21.7  22.2
 22.8  20.   20.6  21.  -17.8  10.9  22.  -11.1  -9.4]
```

```
In [30]: # check altimeter unique values
print(raw_eda_df['altimeter_hpa'].unique())
```


```
[29.82 29.83 29.84 29.86 29.88 29.85 29.9 29.91 29.97 29.98 30.01 30.03
30.05 30.02 30.04 30.07 30.08 30.09 30.06 30.1 30.11 30.12 30.14 30.15
30.16 30.17 30.13 30.19 30.21 30.23 30.24 30.25 30.22 30.2 30.18 30.28
30. 29.99 29.95 29.94 29.93 29.92 29.89 29.96 29.87 29.79 29.8 29.81
30.26 30.27 30.3 30.31 30.33 30.34 30.29 30.35 30.36 30.32 29.78 29.77
29.76 29.61 29.58 29.56 29.54 29.51 29.49 29.46 29.45 29.47 29.48 29.52
29.59 29.63 29.64 29.71 29.72 29.75 29.74 29.73 29.7 29.68 29.67 29.69]
```

In [31]: *# develop a windrose visualization to visualize the most common direction and speeds.*

```
from windrose import WindroseAxes

wind_rose_df = raw_eda_df[['wind_dir_degrees', 'wind_speed_kt']]

ax = WindroseAxes.from_ax()
ax.bar(
    wind_rose_df['wind_dir_degrees'],
    wind_rose_df['wind_speed_kt'],
    normed=True,
    opening=0.8,
    edgecolor='white'
)
ax.set_title('Wind Rose - San Diego Airport')
ax.set_legend(title="Wind Speed (kt)")
plt.show()
```




Destination Exploration

```
In [32]: # check the unique values for DEST
print(raw_eda_df['DEST'].unique())
```

```
[ 'DTW' 'SLC' 'CLT' 'ATL' 'MSP' 'LAS' 'PHX' 'DFW' 'DEN' 'SMF' 'SJC' 'OAK'
  'LAX' 'SFO' 'ORD' 'PHL' 'AUS' 'JFK' 'SEA' 'EWR' 'DAL' 'HNL' 'LIH' 'IAH'
  'MCO' 'BNA' 'IAD' 'OGG' 'BOS' 'RDM' 'MDW' 'ELP' 'RNO' 'HOU' 'HDN' 'KOA'
  'SAT' 'MRY' 'MCI' 'FLL' 'ABQ' 'BOI' 'PDX' 'SBP' 'TUS' 'STS' 'BWI' 'STL'
  'FAT' 'GEG' 'MIA' 'MSY' 'COS' 'DSM' 'PVU' 'BZN' 'JAC' 'EUG' 'FSD' 'BLI'
  'AZA' 'PSC' 'IDA' 'CLE' 'MFR' 'IND' 'FCA' 'MSO' 'TPA' 'PAE' 'ANC' 'PIT'
  'CMH' 'MKE' 'EGE']
```

```
In [33]: # explore visualization of flights by destination
plt.figure(figsize=(12, 6))
sns.countplot(data=raw_eda_df, x='DEST', order=raw_eda_df['DEST'].value_counts().index)
plt.title('Count of Flights by Destination')
plt.xlabel('Destination Airport')
plt.ylabel('Number of Flights')
plt.xticks(rotation=90)
plt.tight_layout()
plt.show()
```



```

In [34]: # map airport codes to region
region_map = {
    # West
    'LAX': 'West', 'SFO': 'West', 'OAK': 'West', 'SJC': 'West', 'SMF': 'West',
    'SBP': 'West', 'STS': 'West', 'SAN': 'West', 'FAT': 'West', 'MRY': 'West',
    'BOI': 'West', 'PDX': 'West', 'RDM': 'West', 'EUG': 'West', 'BLI': 'West',
    'PAE': 'West', 'RNO': 'West', 'PSC': 'West', 'SMF': 'West', 'SEA': 'West',
    'GEG': 'West', 'MFR': 'West',

    # Mountain
    'SLC': 'Mountain', 'DEN': 'Mountain', 'COS': 'Mountain', 'HDN': 'Mountain',
    'ABQ': 'Mountain', 'BZN': 'Mountain', 'JAC': 'Mountain', 'IDA': 'Mountain',
    'MSO': 'Mountain', 'FCA': 'Mountain', 'PVU': 'Mountain', 'EGE': 'Mountain',

    # Southwest
    'PHX': 'Southwest', 'TUS': 'Southwest', 'LAS': 'Southwest', 'AZA': 'Southwest',
    'ELP': 'Southwest',

    # Midwest
    'ORD': 'Midwest', 'MDW': 'Midwest', 'MCI': 'Midwest', 'STL': 'Midwest',
    'DTW': 'Midwest', 'MSP': 'Midwest', 'IND': 'Midwest', 'CMH': 'Midwest',
    'CLE': 'Midwest', 'MKE': 'Midwest', 'DSM': 'Midwest', 'FSD': 'Midwest',

    # South
    'CLT': 'South', 'ATL': 'South', 'BNA': 'South', 'DFW': 'South', 'DAL': 'South', 'AUS': 'South',
    'IAH': 'South', 'HOU': 'South', 'SAT': 'South', 'MSY': 'South', 'MCO': 'South',
    'TPA': 'South', 'FLL': 'South', 'MIA': 'South',

    # Northeast
    'BOS': 'Northeast', 'PHL': 'Northeast', 'JFK': 'Northeast', 'EWR': 'Northeast',
    'PIT': 'Northeast', 'IAD': 'Northeast', 'BWI': 'Northeast',

    # OCONUS will serve as the placeholder for Alaska and Hawaii flights
    # OCONUS refers to Outside Continental United States
    'ANC': 'OCONUS', 'HNL': 'OCONUS', 'KOA': 'OCONUS', 'OGG': 'OCONUS', 'LIH': 'OCONUS',
}
# DEST to region
raw_eda_df['DEST_REGION'] = raw_eda_df['DEST'].map(region_map).fillna('Other')


# results
print(raw_eda_df[['DEST', 'DEST_REGION']].head())

```

	DEST	DEST_REGION
0	DTW	Midwest
1	SLC	Mountain
2	CLT	South
3	ATL	South
4	MSP	Midwest

```
In [35]: # visualize the distribution of flights by region
plt.figure(figsize=(10,6))
sns.countplot(data=raw_eda_df, x='DEST_REGION', order=raw_eda_df['DEST_REGION'].value_counts().index)
plt.title('Count of Flights by Destination Region')
plt.xlabel('Destination Region')
plt.ylabel('Number of Flights')
plt.tight_layout()
plt.show()

# print counts of flights by region
print(raw_eda_df['DEST_REGION'].value_counts())
```



```
DEST_REGION
West      69248
South     33648
Southwest 26551
Mountain  17746
Northeast 17653
Midwest   15044
OCONUS    5710
Name: count, dtype: int64
```

```
In [36]: # creating a stacked bar chart of destination region and flight status
region_status_counts = raw_eda_df.groupby(['DEST_REGION', 'Flight_Status']).size().reset_index()
region_status_counts.columns = ['Region', 'Flight_Status', 'Count']
```



```


# pivot table for plotting
region_status_pivot = region_status_counts.pivot(index='Region', columns='Flight_Status', values='Count').fillna(0)
print(region_status_pivot)

# sort by total flight count per region
region_status_pivot['Total'] = region_status_pivot.sum(axis=1)
region_status_pivot = region_status_pivot.sort_values(by='Total', ascending=False)
region_status_pivot = region_status_pivot.drop(columns='Total')

# plot stacked bar chart
region_status_pivot.plot(kind='bar', stacked=True, figsize=(12, 6))
plt.title('Count of Flights by Destination Region and Flight Status')
plt.xlabel('Destination Region')
plt.ylabel('Number of Flights')
plt.legend(title='Flight Status')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

```

Flight_Status	Cancelled	Delayed	On-Time
Region			
Midwest	128	1115	13801
Mountain	198	1658	15890
Northeast	223	1502	15928
OCONUS	56	609	5045
South	443	3315	29890
Southwest	387	3080	23084
West	683	6392	62173



Date Exploration

In []:

```
In [37]: # One hot encode the DEST_REGION column
dest_region_dummies = pd.get_dummies(raw_eda_df['DEST_REGION'], prefix='Dest_Region')

# concatenate the carrier dummies with the original dataframe, we are not going to get rid of the extra column just yet.
raw_eda_df = pd.concat([raw_eda_df, dest_region_dummies], axis=1)

In [38]: print(raw_eda_df.dtypes)
```

FL_DATE	object
CRS_DEP_TIME	object
DEP_TIME	object
OP_UNIQUE_CARRIER	object
OP_CARRIER_FL_NUM	int64
DEST	object
DEP_DEL15	int64
CANCELLED	int64
CANCELLATION_CODE	object
CARRIER_DELAY	int64
WEATHER_DELAY	int64
NAS_DELAY	int64
SECURITY_DELAY	int64
LATE_AIRCRAFT_DELAY	int64
wind_dir_degrees	int64
wind_speed_kt	int64
wind_gust_kt	int64
visibility_statute_mi	float64
temperature_c	float64
dewpoint_c	float64
altimeter_hpa	float64
Flight_Status	object
Carrier_AA	bool
Carrier_AS	bool
Carrier_B6	bool
Carrier_DL	bool
Carrier_F9	bool
Carrier_G4	bool
Carrier_HA	bool
Carrier_NK	bool
Carrier_OO	bool
Carrier_UA	bool
Carrier_WN	bool
DEST_REGION	object
Dest_Region_Midwest	bool
Dest_Region_Mountain	bool
Dest_Region_Northeast	bool
Dest_Region_OCONUS	bool
Dest_Region_South	bool
Dest_Region_Southwest	bool
Dest_Region_West	bool

dtype: object

Finalize data frame for processed data

```
In [39]: processed_df = raw_eda_df.copy()
processed_df = processed_df.drop(columns=['DEP_TIME', 'DEP_DEL15', 'CANCELLED', 'OP_CARRIER_FL_NUM', 'CANCELLATION_CODE', 'CARRI

# create the target variable for the model using Flight_Status where if Cancelled or Delayed = 1 and On-time = 0
processed_df['Flight_Status_Binary'] = processed_df['Flight_Status'].apply(
    lambda x: 1 if x in ['Delayed', 'Cancelled'] else 0
)

#print head
print(processed_df.head())
```

	FL_DATE	CRS_DEP_TIME	OP_UNIQUE_CARRIER	DEST	wind_dir_degrees	\
0	2023-01-01	06:10:00		DL DTW	290	
1	2023-01-01	06:15:00		DL SLC	290	
2	2023-01-01	06:15:00	AA	CLT	290	
3	2023-01-01	06:15:00	DL	ATL	290	
4	2023-01-01	06:15:00	DL	MSP	290	

	wind_speed_kt	wind_gust_kt	visibility_statute_mi	temperature_c	\
0	7	19	6.0	13.3	
1	7	19	6.0	13.3	
2	7	19	6.0	13.3	
3	7	19	6.0	13.3	
4	7	19	6.0	13.3	

	dewpoint_c	...	Carrier_WN	DEST_REGION	Dest_Region_Midwest	\
0	10.0	...	False	Midwest	True	
1	10.0	...	False	Mountain	False	
2	10.0	...	False	South	False	
3	10.0	...	False	South	False	
4	10.0	...	False	Midwest	True	

	Dest_Region_Mountain	Dest_Region_Northeast	Dest_Region_OCONUS	\
0	False	False	False	
1	True	False	False	
2	False	False	False	
3	False	False	False	
4	False	False	False	

	Dest_Region_South	Dest_Region_Southwest	Dest_Region_West	\
0	False	False	False	
1	False	False	False	
2	True	False	False	
3	True	False	False	
4	False	False	False	

	Flight_Status_Binary
0	0
1	0
2	1
3	0
4	0

[5 rows x 32 columns]

```
In [40]: # save the processed data to a csv file
```

```
processed_df.to_csv("processed_data/processed_data.csv", index=False)
```

Upload processed_data.csv to Processed_data S3 Bucket

```
In [41]: # examine list of buckets available
s3 = boto3.client('s3')
response = s3.list_buckets()
print("Existing buckets:")
for bucket in response['Buckets']:
    print(bucket['Name'])
```

Existing buckets:

```
group9-ml-proj-athena-query-bucket-grw
group9-ml-proj-model-artifacts-bucket-grw
group9-ml-proj-monitoring-bucket-grw
group9-ml-proj-processed-data-bucket-grw
group9-ml-proj-raw-data-bucket-grw
group9-ml-proj-training-artifacts-bucket-grw
sagemaker-studio-7sb8gxpljq9
sagemaker-studio-chxawrkbs1q
sagemaker-studio-o4i4ukj5y3b
sagemaker-us-east-1-804823187130
```

```
In [42]: # set local path of organized weather CSV file
local_file_path = 'processed_data/processed_data.csv'
```

```
In [43]: # set the raw data bucket variable
processed_bucket = "group9-ml-proj-processed-data-bucket-grw"
```

```
In [44]: # set the S3 key path, there will be
s3_key = "all_data/processed_data.csv"
```

```
In [45]: # upload file to the S3 bucket
try:
    s3.upload_file(local_file_path, processed_bucket, s3_key)
    print(f"Uploaded {local_file_path} to s3://{processed_bucket}/{s3_key}")
except Exception as e:
    print("Error uploading file to S3:", e)
```

Uploaded processed_data/processed_data.csv to s3://group9-ml-proj-processed-data-bucket-grw/all_data/processed_data.csv

```
In [46]: %%html
<p><b>Shutting down your kernel for this notebook to release resources.</b></p>
<button class="sm-command-button" data-commandlinker-command="kernelmenu:shutdown" style="display:none;">Shutdown Kernel</button>
```

```
<script>
try {
  els = document.getElementsByClassName("sm-command-button");
  els[0].click();
}
catch(err) {
  // NoOp
}
</script>
```

Shutting down your kernel for this notebook to release resources.

In []:

```
In [1]: !pip install imblearn
```

```
Collecting imblearn
  Downloading imblearn-0.0-py2.py3-none-any.whl.metadata (355 bytes)
Collecting imbalanced-learn (from imblearn)
  Downloading imbalanced_learn-0.13.0-py3-none-any.whl.metadata (8.8 kB)
Requirement already satisfied: numpy<3,>=1.24.3 in /opt/conda/lib/python3.11/site-packages (from imbalanced-learn->imblearn) (1.26.4)
Requirement already satisfied: scipy<2,>=1.10.1 in /opt/conda/lib/python3.11/site-packages (from imbalanced-learn->imblearn) (1.15.2)
Requirement already satisfied: scikit-learn<2,>=1.3.2 in /opt/conda/lib/python3.11/site-packages (from imbalanced-learn->imblearn) (1.5.2)
Collecting sklearn-compat<1,>=0.1 (from imbalanced-learn->imblearn)
  Downloading sklearn_compat-0.1.3-py3-none-any.whl.metadata (18 kB)
Requirement already satisfied: joblib<2,>=1.1.1 in /opt/conda/lib/python3.11/site-packages (from imbalanced-learn->imblearn) (1.4.2)
Requirement already satisfied: threadpoolctl<4,>=2.0.0 in /opt/conda/lib/python3.11/site-packages (from imbalanced-learn->imblearn) (3.5.0)
Downloading imblearn-0.0-py2.py3-none-any.whl (1.9 kB)
Downloading imbalanced_learn-0.13.0-py3-none-any.whl (238 kB)
Downloading sklearn_compat-0.1.3-py3-none-any.whl (18 kB)
Installing collected packages: sklearn-compat, imbalanced-learn, imblearn
Successfully installed imbalanced-learn-0.13.0 imblearn-0.0 sklearn-compat-0.1.3
```

```
In [1]: import pandas as pd
import warnings
warnings.filterwarnings("ignore")
import boto3
import time
import io
```

```
In [2]: # Athena Query of the S3 bucket function
def run_athena_query(query, database, output_location):
    athena_client = boto3.client('athena')

    # start the query execution.
    response = athena_client.start_query_execution(
        QueryString=query,
        QueryExecutionContext={'Database': database},
        ResultConfiguration={'OutputLocation': output_location}
    )
    query_execution_id = response['QueryExecutionId']
    print(f"Query submitted. Execution ID: {query_execution_id}")
```



```

# poll the query status until it completes.
state = 'RUNNING'
while state in ['RUNNING', 'QUEUED']:
    response = athena_client.get_query_execution(QueryExecutionId=query_execution_id)
    state = response['QueryExecution']['Status']['State']
    print(f"Current query state: {state}")
    if state in ['RUNNING', 'QUEUED']:
        time.sleep(2)

if state == 'FAILED':
    reason = response['QueryExecution']['Status'].get('StateChangeReason', 'Unknown error')
    raise Exception(f"Query {query_execution_id} failed: {reason}")

print(f"Query {query_execution_id} completed successfully.")
return query_execution_id

```

```

In [3]: # turn query results into DF function
def get_query_results_as_dataframe(query_execution_id):
    athena_client = boto3.client('athena')
    s3_client = boto3.client('s3')

    # get query execution to find the S3 location of the results.
    response = athena_client.get_query_execution(QueryExecutionId=query_execution_id)
    result_location = response['QueryExecution']['ResultConfiguration']['OutputLocation']

    # parse S3 bucket and key from the result loc.
    s3_path = result_location.replace("s3://", "").split("/", 1)
    bucket = s3_path[0]
    key = s3_path[1]

    # get the CSV result file from S3.
    obj = s3_client.get_object(Bucket=bucket, Key=key)
    df = pd.read_csv(io.BytesIO(obj['Body'].read()))
    return df

```

```

In [5]: # configuration and Query Execution
# set Athena db name and output location.
athena_database = 'processed_data'
output_location = 's3://group9-ml-proj-athena-query-bucket-grw/'

# define your query.
query = """
SELECT wind_dir_degrees,
       wind_speed_kt,
       wind_gust_kt,

```

```

        visibility_statute_mi,
        temperature_c,
        dewpoint_c,
        altimeter_hpa,
        Carrier_AA,
        Carrier_AS,
        Carrier_B6,
        Carrier_DL,
        Carrier_F9,
        Carrier_G4,
        Carrier_HA,
        Carrier_NK,
        Carrier_OO,
        Carrier_UA,
        Carrier_WN,
        Dest_Region_Midwest,
        Dest_Region_Mountain,
        Dest_Region_Northeast,
        Dest_Region_OCONUS,
        Dest_Region_South,
        Dest_Region_Southwest,
        Dest_Region_West,
        Flight_Status_Binary
FROM processed_data;
"""

# run the Athena query and return execution ID.
query_execution_id = run_athena_query(query, database=athena_database, output_location=output_location)

# retrieve the query results as a pandas df named 'data'
data = get_query_results_as_dataframe(query_execution_id)

# display the first few rows of the data
print(data.head())

```

Query submitted. Execution ID: 060667fd-9695-49e3-bba8-d822127a50a0

Current query state: QUEUED

Current query state: RUNNING

Current query state: SUCCEEDED

Query 060667fd-9695-49e3-bba8-d822127a50a0 completed successfully.

	wind_dir_degrees	wind_speed_kt	wind_gust_kt	visibility_statute_mi	\
0	290.0	7.0	19.0	6.0	
1	290.0	7.0	19.0	6.0	
2	290.0	7.0	19.0	6.0	
3	290.0	7.0	19.0	6.0	
4	290.0	7.0	19.0	6.0	

	temperature_c	dewpoint_c	altimeter_hpa	Carrier_AA	Carrier_AS	\
0	13.3	10.0	29.82	False	False	
1	13.3	10.0	29.82	False	False	
2	13.3	10.0	29.82	True	False	
3	13.3	10.0	29.82	False	False	
4	13.3	10.0	29.82	False	False	

	Carrier_B6	...	Carrier_UA	Carrier_WN	Dest_Region_Midwest	\
0	False	...	False	False	True	
1	False	...	False	False	False	
2	False	...	False	False	False	
3	False	...	False	False	False	
4	False	...	False	False	True	

	Dest_Region_Mountain	Dest_Region_Northeast	Dest_Region_OCONUS	\
0	False	False	False	
1	True	False	False	
2	False	False	False	
3	False	False	False	
4	False	False	False	

	Dest_Region_South	Dest_Region_Southwest	Dest_Region_West	\
0	False	False	False	
1	False	False	False	
2	True	False	False	
3	True	False	False	
4	False	False	False	

	Flight_Status_Binary
0	False
1	False
2	False
3	False
4	False

[5 rows x 26 columns]

Split 70/15/15

Undersample

```
In [4]: from sklearn.model_selection import train_test_split
        from imblearn.under_sampling import RandomUnderSampler

        # Define features (X) and target (y)
        X = data.drop(['Flight_Status_Binary'], axis=1)
        y = data['Flight_Status_Binary']

        # Step 1: Split into training (70%) and temp dataset (30%)
        X_train, X_temp, y_train, y_temp = train_test_split(
            X, y, test_size=0.30, stratify=y, random_state=42
        )

        # Step 2: Further split temp into validation (15%) and testing (15%)
        X_val, X_test, y_val, y_test = train_test_split(
            X_temp, y_temp, test_size=0.50, stratify=y_temp, random_state=42
        )

        # Step 3: Perform undersampling on the training dataset only
        rus = RandomUnderSampler(random_state=42)
        X_train_resampled, y_train_resampled = rus.fit_resample(X_train, y_train)

        # Verify the shape and class distributions
        print("Dataset Shapes:")
        print(f"Original Training Set: {X_train.shape}")
        print(f"Resampled Training Set: {X_train_resampled.shape}")
        print(f"Validation Set: {X_val.shape}")
        print(f"Testing Set: {X_test.shape}")

        print("\nTraining Set Class Distribution (Original):")
        print(y_train.value_counts())

        print("\nTraining Set Class Distribution (After Resampling):")
        print(y_train_resampled.value_counts())
```

Dataset Shapes:

Original Training Set: (129920, 25)

Resampled Training Set: (27704, 25)

Validation Set: (27840, 25)

Testing Set: (27840, 25)

Training Set Class Distribution (Original):

Flight_Status_Binary

0 116068

1 13852

Name: count, dtype: int64

Training Set Class Distribution (After Resampling):

Flight_Status_Binary

0 13852

1 13852

Name: count, dtype: int64

Modeling Building

Logistic Model

```
In [5]: from sklearn.linear_model import LogisticRegression
        from sklearn.metrics import classification_report, confusion_matrix

        # Instantiate the model
        model = LogisticRegression(max_iter=1000, random_state=42)

        # Train model on resampled training data
        model.fit(X_train_resampled, y_train_resampled)

        # Evaluate on validation data
        y_pred_val = model.predict(X_val)

        # Performance evaluation
        print(classification_report(y_val, y_pred_val))
        print(confusion_matrix(y_val, y_pred_val))
```

	precision	recall	f1-score	support
0	0.92	0.57	0.70	24871
1	0.14	0.59	0.23	2969
accuracy			0.57	27840
macro avg	0.53	0.58	0.46	27840
weighted avg	0.84	0.57	0.65	27840

```
[[14123 10748]
 [ 1217 1752]]
```

Random Forest

```
In [6]: from sklearn.ensemble import RandomForestClassifier

rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train_resampled, y_train_resampled)

y_val_pred_rf = rf.predict(X_val)
print("Random Forest Validation Report:")
print(classification_report(y_val, y_val_pred_rf))
```

Random Forest Validation Report:

	precision	recall	f1-score	support
0	0.94	0.64	0.76	24871
1	0.18	0.66	0.28	2969
accuracy			0.65	27840
macro avg	0.56	0.65	0.52	27840
weighted avg	0.86	0.65	0.71	27840

Xgboost

```
In [7]: from xgboost import XGBClassifier

xgb = XGBClassifier(use_label_encoder=False, eval_metric='logloss', random_state=42)
xgb.fit(X_train_resampled, y_train_resampled)

y_val_pred_xgb = xgb.predict(X_val)
```

```
print("XGBoost Validation Report:")
print(classification_report(y_val, y_val_pred_xgb))
```

XGBoost Validation Report:

	precision	recall	f1-score	support
0	0.95	0.66	0.78	24871
1	0.19	0.69	0.30	2969
accuracy			0.66	27840
macro avg	0.57	0.67	0.54	27840
weighted avg	0.87	0.66	0.73	27840

Since our goal number 2 is to reduce delay times by informing logistics, its better to catch more delays. Using a lower threshold to increase recall, since missing a delay would be worse than predicting one that doesn't end up happening

Find which is the best threshold with AUC

```
In [8]: from sklearn.ensemble import RandomForestClassifier

# Train a simple random forest model on data
best_rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
best_rf_model.fit(X_train_resampled, y_train_resampled)
```

```
Out[8]: RandomForestClassifier
RandomForestClassifier(random_state=42)
```


```
In [9]: # Predict probabilities for the positive class (delays/cancellations)
y_proba_rf = best_rf_model.predict_proba(X_test)[ :, 1]
```

```
In [10]: from sklearn.metrics import precision_recall_curve, auc
import matplotlib.pyplot as plt

precision, recall, thresholds = precision_recall_curve(y_test, y_proba_rf)
pr_auc = auc(recall, precision)

plt.figure(figsize=(8,6))
plt.plot(recall, precision, label=f'PR Curve (AUC={pr_auc:.2f})')
```

```
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.legend()
plt.grid(True)
plt.show()
```



```
In [11]: import numpy as np
import pandas as pd
from sklearn.metrics import precision_score, recall_score, f1_score

thresholds = np.arange(0.1, 0.51, 0.05)
results = []
```



```

for threshold in thresholds:
    y_pred_threshold = (y_proba_rf >= threshold).astype(int)
    precision = precision_score(y_test, y_pred_threshold)
    recall = recall_score(y_test, y_pred_threshold)
    f1 = f1_score(y_test, y_pred_threshold)
    results.append([threshold, precision, recall, f1])

# Convert results to DataFrame
results_df = pd.DataFrame(results, columns=['Threshold', 'Precision', 'Recall', 'F1-score'])
print(results_df)

```

	Threshold	Precision	Recall	F1-score
0	0.10	0.111495	0.990229	0.200423
1	0.15	0.114386	0.969003	0.204617
2	0.20	0.118907	0.950135	0.211363
3	0.25	0.124584	0.920148	0.219454
4	0.30	0.133171	0.884771	0.231498
5	0.35	0.142210	0.835916	0.243068
6	0.40	0.154194	0.792790	0.258174
7	0.45	0.164611	0.719340	0.267913
8	0.50	0.181361	0.655660	0.284129

```

In [12]: final_threshold = 0.25
y_pred_final = (y_proba_rf >= final_threshold).astype(int)

print("Final Threshold (0.25) Classification Report:")
print(classification_report(y_test, y_pred_final))
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred_final))

```

Final Threshold (0.25) Classification Report:

	precision	recall	f1-score	support
0	0.96	0.22	0.36	24872
1	0.12	0.92	0.22	2968
accuracy			0.30	27840
macro avg	0.54	0.57	0.29	27840
weighted avg	0.87	0.30	0.35	27840

Confusion Matrix:

```

[[ 5511 19361]
 [  226  2742]]

```

Random Forest with 0.25

```
In [13]: # Predict probabilities
y_proba_rf = best_rf_model.predict_proba(X_test)[: , 1]

# Set optimal threshold
threshold = 0.25
y_pred_rf = (y_proba_rf >= threshold).astype(int)

# Evaluation metrics
print("Random Forest with threshold = 0.25")
print(classification_report(y_test, y_pred_rf))
print(confusion_matrix(y_test, y_pred_rf))
```

```
Random Forest with threshold = 0.25
              precision    recall  f1-score   support

      0       0.96      0.22      0.36      24872
      1       0.12      0.92      0.22       2968

 accuracy      0.30      0.30      0.30      27840
 macro avg       0.54      0.57      0.29      27840
weighted avg       0.87      0.30      0.35      27840

[[ 5511 19361]
 [  226  2742]]
```

Logistic with 0.25

```
In [14]: from sklearn.linear_model import LogisticRegression

# Train Logistic Regression on resampled training data
lr_model = LogisticRegression(max_iter=1000, random_state=42)
lr_model.fit(X_train_resampled, y_train_resampled)

# Predict probabilities
y_proba_lr = lr_model.predict_proba(X_test)[: , 1]

# Apply optimal threshold
y_pred_lr = (y_proba_lr >= threshold).astype(int)

# Print
```

```
print("Logistic Regression with threshold = 0.25")
print(classification_report(y_test, y_pred_lr))
print(confusion_matrix(y_test, y_pred_lr))
```

```
Logistic Regression with threshold = 0.25
```

	precision	recall	f1-score	support
0	0.97	0.01	0.02	24872
1	0.11	1.00	0.19	2968
accuracy			0.11	27840
macro avg	0.54	0.50	0.11	27840
weighted avg	0.88	0.11	0.04	27840

```
[[ 213 24659]
 [   6  2962]]
```

XGBoost with 0.25 threshold

In [15]: `from xgboost import XGBClassifier`

```
# Train XGBoost model
xgb_model = XGBClassifier(use_label_encoder=False, eval_metric='logloss', random_state=42)
xgb_model.fit(X_train_resampled, y_train_resampled)

# Predict probabilities
y_proba_xgb = xgb_model.predict_proba(X_test)[:, 1]

# Apply optimal threshold
y_pred_xgb = (y_proba_xgb >= threshold).astype(int)

# Evaluation metrics
print("XGBoost with threshold = 0.25")
print(classification_report(y_test, y_pred_xgb))
print(confusion_matrix(y_test, y_pred_xgb))
```

```
XGBoost with threshold = 0.25
```

	precision	recall	f1-score	support
0	0.97	0.23	0.37	24872
1	0.13	0.95	0.23	2968
accuracy			0.31	27840
macro avg	0.55	0.59	0.30	27840
weighted avg	0.88	0.31	0.36	27840

```
[[ 5739 19133]
 [ 158 2810]]
```

Best model:

After evaluating all three models, Random Forest is the best option for our goal. It catches the most delays (92% recall), which is crucial because missing a delay costs more than occasionally predicting a false alarm. Even though precision is lower, this trade-off makes sense for our business scenario. Let's move forward with the Random Forest model at a 0.25 threshold to best support logistics planning and minimize disruptions.

Random Forest (Hyperparameters):

```
In [16]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import RandomizedSearchCV

# Hyperparameter grid for tuning
param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [10, 20, None],
    'min_samples_split': [2, 5, 10],
    'class_weight': [None, 'balanced']
}

# Randomized search for best hyperparameters
rf_tuned = RandomizedSearchCV(
    RandomForestClassifier(random_state=42),
    param_distributions=param_grid,
    scoring='recall', # prioritize recall
    cv=5,
    n_iter=5,
    random_state=42,
    verbose=1
)
```

```
# Fit the tuned Random Forest model
rf_tuned.fit(X_train_resampled, y_train_resampled)

# Check best hyperparameters
print("Best Hyperparameters:", rf_tuned.best_params_)
```

Fitting 5 folds for each of 5 candidates, totalling 25 fits

Best Hyperparameters: {'n_estimators': 100, 'min_samples_split': 5, 'max_depth': 20, 'class_weight': None}

```
In [17]: # Final best parameters
final_rf_model = rf_tuned.best_estimator_

# Predict probabilities on test data
y_proba_final_rf = final_rf_model.predict_proba(X_test)[: , 1]

# Apply threshold (0.25)
final_threshold = 0.25
y_pred_final_rf = (y_proba_final_rf >= final_threshold).astype(int)

# Final evaluation
print("Final Random Forest Model (Tuned Hyperparameters, Threshold = 0.25):")
print(classification_report(y_test, y_pred_final_rf))
print(confusion_matrix(y_test, y_pred_final_rf))
```

Final Random Forest Model (Tuned Hyperparameters, Threshold = 0.25):

	precision	recall	f1-score	support
0	0.98	0.11	0.20	24872
1	0.12	0.98	0.21	2968
accuracy			0.21	27840
macro avg	0.55	0.54	0.21	27840
weighted avg	0.88	0.21	0.20	27840

```
[[ 2825 22047]
 [   71  2897]]
```

Conclusion:

After tuning the hyperparameters for our Random Forest model and applying our optimal threshold (0.25), we achieved a good performance with 98% of actual delays. Although this high recall resulted in some false alarms (low precision of 12%), this trade-off is acceptable since missing real delays would cost our logistics operation far more than a few unnecessary alerts.

Training model to cloud with Sagemaker

1. Converted local scikit-learn model into a SageMaker training script (`train.py`)
2. Created `train.csv` and `validation.csv` from our 70/15/15 split
3. Uploaded the datasets to S3 for SageMaker to access
4. Added missing dependency `imblearn` via a `requirements.txt` file
5. Used `sagemaker.sklearn.estimator.SKLearn` to launch the job in the cloud
6. Model trained successfully using a managed `m1.m5.xlarge` instance
7. Model output automatically saved to S3 — fully cloud-compliant

```
In [18]: # Save training and validation sets to CSV
train_df = X_train.copy()
train_df["Flight_Status_Binary"] = y_train
train_df.to_csv("train.csv", index=False)

val_df = X_val.copy()
val_df["Flight_Status_Binary"] = y_val
val_df.to_csv("validation.csv", index=False)
```

```
In [19]: train_input = session.upload_data("Model_data/train.csv", key_prefix="ads-508/train")
val_input = session.upload_data("Model_data/validation.csv", key_prefix="ads-508/validation")
```

```
-----
NameError                                Traceback (most recent call last)
Cell In[19], line 1
----> 1 train_input = session.upload_data("Model_data/train.csv", key_prefix="ads-508/train")
      2 val_input = session.upload_data("Model_data/validation.csv", key_prefix="ads-508/validation")

NameError: name 'session' is not defined
```

```
In [ ]: import sagemaker
from sagemaker import get_execution_role
from sagemaker.sklearn.estimator import SKLearn

# Set up session and role
session = sagemaker.Session()
role = get_execution_role()

# Define the estimator
sklearn_estimator = SKLearn(
```

```

    entry_point="Model_data/train.py",
    role=role,
    instance_type= "ml.m5.large",
    framework_version="1.0-1",
    py_version="py3",
    sagemaker_session=session,
    dependencies=["Model_data/requirements.txt"]
)

# Launch SageMaker training job
sklearn_estimator.fit({
    "train": train_input,
    "validation": val_input
})

```

```

In [ ]: from sagemaker.sklearn.estimator import SKLearn

sklearn_estimator = SKLearn.attach("sagemaker-scikit-learn-2025-03-29-21-43-12-337")

```

```

In [ ]: predictor = sklearn_estimator.deploy(
    initial_instance_count=1,
    instance_type="ml.m5.large"
)

```

Comments for Professor:

Errors Encountered During Development-->

ResourceLimitExceeded:

The requested resource ml.t3.medium is not available in this region This occurred when trying to launch a training job using the ml.t3.medium instance type. The instance was not available in the us-east-1 region, leading to a resource limit error.

AccessDeniedException:

Not authorized to perform: sagemaker: CreateModel After switching to an available instance type (ml.m5.xlarge and ml.m5.large), an error occurred when attempting to deploy the trained model as an inference endpoint. This was due to an IAM policy restriction that explicitly denied the sagemaker:CreateModel action. Since deployment was not required for this phase of the assignment, the project proceeded without creating an endpoint.

Final Model built-in SageMaker XGBoost Estimator

```
In [21]: import sagemaker
        from sagemaker import get_execution_role

        # Set up SageMaker session and role
        session = sagemaker.Session()
        role = get_execution_role()
```

```
In [23]: # Upload data to S3
        train_input = session.upload_data("Model_data/train.csv", key_prefix="ads-508/train")
        val_input = session.upload_data("Model_data/validation.csv", key_prefix="ads-508/validation")
```

```
In [25]: from sagemaker.inputs import TrainingInput
        from sagemaker.xgboost.estimator import XGBoost

        xgb_estimator = XGBoost(
            entry_point="Model_data/xgboost_train.py",
            role=role,
            instance_type="ml.m5.large",
            instance_count=1,
            framework_version="1.3-1",
            py_version="py3",
            sagemaker_session=session
        )

        xgb_estimator.fit({
            "train": TrainingInput(train_input, content_type="csv"),
            "validation": TrainingInput(val_input, content_type="csv")
        })
```

[04/02/25 04:27:26] INFO Ignoring unnecessary Python version: py3.

image_uris.py:603

INFO Ignoring unnecessary instance type: ml.m5.large.

image_uris.py:530

INFO SageMaker Python SDK will collect telemetry to help us better understand our user's needs, diagnose issues, and deliver additional features. telemetry_logging.py:91
To opt out of telemetry, please disable via TelemetryOptOut parameter in SDK defaults config. For more information, refer to <https://sagemaker.readthedocs.io/en/stable/overview.html#configuring-and-using-defaults-with-the-sagemaker-python-sdk>.

INFO Creating training-job with name: session.py:1042
sagemaker-xgboost-2025-04-02-04-27-26-344

```

2025-04-02 04:27:28 Starting - Starting the training job...
..25-04-02 04:27:42 Starting - Preparing the instances for training.
..25-04-02 04:28:06 Downloading - Downloading input data.
..25-04-02 04:28:51 Downloading - Downloading the training image.
.[2025-04-02 04:29:31.953 ip-10-2-76-153.ec2.internal:7 INFO utils.py:28] RULE_JOB_STOP_SIGNAL_FILENAME: None
[2025-04-02 04:29:31.977 ip-10-2-76-153.ec2.internal:7 INFO profiler_config_parser.py:111] User has disabled profiler.
[2025-04-02:04:29:31:INFO] Imported framework sagemaker_xgboost_container.training
[2025-04-02:04:29:31:INFO] No GPUs detected (normal if no gpus installed)
[2025-04-02:04:29:31:INFO] Invoking user training script.
[2025-04-02:04:29:32:INFO] Module xgboost_train does not provide a setup.py.
Generating setup.py
[2025-04-02:04:29:32:INFO] Generating setup.cfg
[2025-04-02:04:29:32:INFO] Generating MANIFEST.in
[2025-04-02:04:29:32:INFO] Installing module with the following command:
/miniconda3/bin/python3 -m pip install .
Processing /opt/ml/code
  Preparing metadata (setup.py): started
  Preparing metadata (setup.py): finished with status 'done'
Building wheels for collected packages: xgboost-train
  Building wheel for xgboost-train (setup.py): started
  Building wheel for xgboost-train (setup.py): finished with status 'done'
  Created wheel for xgboost-train: filename=xgboost_train-1.0.0-py2.py3-none-any.whl size=3626 sha256=d6f76168516b2a6dee8ada06770
f07887413418ed6e0701b30202a18bc67c437
  Stored in directory: /home/model-server/tmp/pip-ephem-wheel-cache-juajazdb/wheels/3e/0f/51/2f1df833dd0412c1bc2f5ee56baac195b5be
563353d111dca6
Successfully built xgboost-train
Installing collected packages: xgboost-train
Successfully installed xgboost-train-1.0.0
WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manage
r. It is recommended to use a virtual environment instead: https://pip.pypa.io/warnings/venv
[notice] A new release of pip is available: 23.0.1 -> 24.0
[notice] To update, run: pip install --upgrade pip
[2025-04-02:04:29:33:INFO] No GPUs detected (normal if no gpus installed)
[2025-04-02:04:29:34:INFO] Invoking user script
Training Env:
{
  "additional_framework_parameters": {},
  "channel_input_dirs": {
    "train": "/opt/ml/input/data/train",
    "validation": "/opt/ml/input/data/validation"
  },
  "current_host": "algo-1",
  "framework_module": "sagemaker_xgboost_container.training:main",
  "hosts": [
    "algo-1"
  ],

```

```
"hyperparameters": {},
"input_config_dir": "/opt/ml/input/config",
"input_data_config": {
    "train": {
        "ContentType": "csv",
        "TrainingInputMode": "File",
        "S3DistributionType": "FullyReplicated",
        "RecordWrapperType": "None"
    },
    "validation": {
        "ContentType": "csv",
        "TrainingInputMode": "File",
        "S3DistributionType": "FullyReplicated",
        "RecordWrapperType": "None"
    }
},
"input_dir": "/opt/ml/input",
"is_master": true,
"job_name": "sagemaker-xgboost-2025-04-02-04-27-26-344",
"log_level": 20,
"master_hostname": "algo-1",
"model_dir": "/opt/ml/model",
"module_dir": "s3://sagemaker-us-east-1-230393891640/sagemaker-xgboost-2025-04-02-04-27-26-344/source/sourcedir.tar.gz",
"module_name": "xgboost_train",
"network_interface_name": "eth0",
"num_cpus": 2,
"num_gpus": 0,
"output_data_dir": "/opt/ml/output/data",
"output_dir": "/opt/ml/output",
"output_intermediate_dir": "/opt/ml/output/intermediate",
"resource_config": {
    "current_host": "algo-1",
    "current_instance_type": "ml.m5.large",
    "current_group_name": "homogeneousCluster",
    "hosts": [
        "algo-1"
    ],
    "instance_groups": [
        {
            "instance_group_name": "homogeneousCluster",
            "instance_type": "ml.m5.large",
            "hosts": [
                "algo-1"
            ]
        }
    ]
},
],
```

```

        "network_interface_name": "eth0"
    },
    "user_entry_point": "xgboost_train.py"
}
Environment variables:
SM_HOSTS=["algo-1"]
SM_NETWORK_INTERFACE_NAME=eth0
SM_HPS={}
SM_USER_ENTRY_POINT=xgboost_train.py
SM_FRAMEWORK_PARAMS={}
SM_RESOURCE_CONFIG={"current_group_name":"homogeneousCluster","current_host":"algo-1","current_instance_type":"ml.m5.large","hosts":["algo-1"],"instance_groups":[{"hosts":["algo-1"],"instance_group_name":"homogeneousCluster","instance_type":"ml.m5.large"}],"network_interface_name":"eth0"}
SM_INPUT_DATA_CONFIG={"train":{"ContentType":"csv","RecordWrapperType":"None","S3DistributionType":"FullyReplicated","TrainingInputMode":"File"},"validation":{"ContentType":"csv","RecordWrapperType":"None","S3DistributionType":"FullyReplicated","TrainingInputMode":"File"}}
SM_OUTPUT_DATA_DIR=/opt/ml/output/data
SM_CHANNELS=["train","validation"]
SM_CURRENT_HOST=algo-1
SM_MODULE_NAME=xgboost_train
SM_LOG_LEVEL=20
SM_FRAMEWORK_MODULE=sagemaker_xgboost_container.training:main
SM_INPUT_DIR=/opt/ml/input
SM_INPUT_CONFIG_DIR=/opt/ml/input/config
SM_OUTPUT_DIR=/opt/ml/output
SM_NUM_CPUS=2
SM_NUM_GPUS=0
SM_MODEL_DIR=/opt/ml/model
SM_MODULE_DIR=s3://sagemaker-us-east-1-230393891640/sagemaker-xgboost-2025-04-02-04-27-26-344/source/sourcedir.tar.gz
SM_TRAINING_ENV={"additional_framework_parameters":{},"channel_input_dirs":{"train":"/opt/ml/input/data/train","validation":"/opt/ml/input/data/validation"},"current_host":"algo-1","framework_module":"sagemaker_xgboost_container.training:main","hosts":["algo-1"],"hyperparameters":{},"input_config_dir":"/opt/ml/input/config","input_data_config":{"train":{"ContentType":"csv","RecordWrapperType":"None","S3DistributionType":"FullyReplicated","TrainingInputMode":"File"},"validation":{"ContentType":"csv","RecordWrapperType":"None","S3DistributionType":"FullyReplicated","TrainingInputMode":"File"}},"input_dir":"/opt/ml/input","is_master":true,"job_name":"sagemaker-xgboost-2025-04-02-04-27-26-344","log_level":20,"master_hostname":"algo-1","model_dir":"/opt/ml/model","module_dir":"s3://sagemaker-us-east-1-230393891640/sagemaker-xgboost-2025-04-02-04-27-26-344/source/sourcedir.tar.gz","module_name":"xgboost_train","network_interface_name":"eth0","num_cpus":2,"num_gpus":0,"output_data_dir":"/opt/ml/output/data","output_dir":"/opt/ml/output","output_intermediate_dir":"/opt/ml/output/intermediate","resource_config":{"current_group_name":"homogeneousCluster","current_host":"algo-1","current_instance_type":"ml.m5.large","hosts":["algo-1"],"instance_groups":[{"hosts":["algo-1"],"instance_group_name":"homogeneousCluster","instance_type":"ml.m5.large"}],"network_interface_name":"eth0"},"user_entry_point":"xgboost_train.py"}
SM_USER_ARGS=[]
SM_OUTPUT_INTERMEDIATE_DIR=/opt/ml/output/intermediate
SM_CHANNEL_TRAIN=/opt/ml/input/data/train
SM_CHANNEL_VALIDATION=/opt/ml/input/data/validation
PYTHONPATH=/miniconda3/bin:/miniconda3/lib/python/site-packages/xgboost/dmlc-core/tracker:/miniconda3/lib/python37.zip:/minicon

```

da3/lib/python3.7:/miniconda3/lib/python3.7/lib-dynload:/miniconda3/lib/python3.7/site-packages
Invoking script with the following command:
/miniconda3/bin/python3 -m xgboost_train

2025-04-02 04:29:52 Uploading - Uploading generated training model

2025-04-02 04:30:11 Completed - Training job completed

Training seconds: 125

Billable seconds: 125

Shutdown Kernel

```
In [1]: %%html
<p><b>Shutting down your kernel for this notebook to release resources.</b></p>
<button class="sm-command-button" data-commandlinker-command="kernelmenu:shutdown" style="display:none;">Shutdown Kernel</button>

<script>
try {
  els = document.getElementsByClassName("sm-command-button");
  els[0].click();
}
catch(err) {
  // NoOp
}
</script>
```

Shutting down your kernel for this notebook to release resources.